

# Building Next-Gen IXP Networks

A Containerlab Workshop



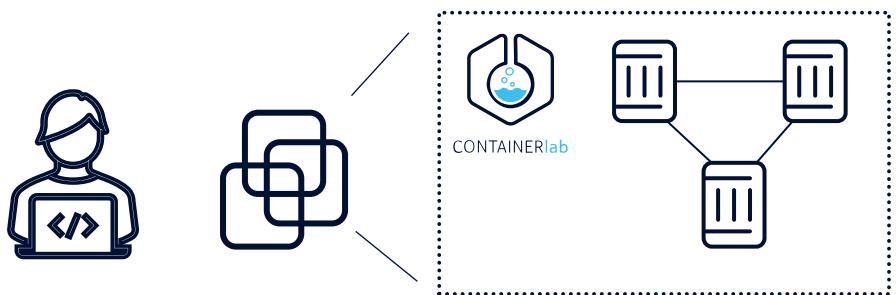
Bastien Claeys  
Paresh Khatri  
Thomas Corre

Workshop repository

Repo & slides

[github.com/thcorre/innog8-workshop](https://github.com/thcorre/innog8-workshop)

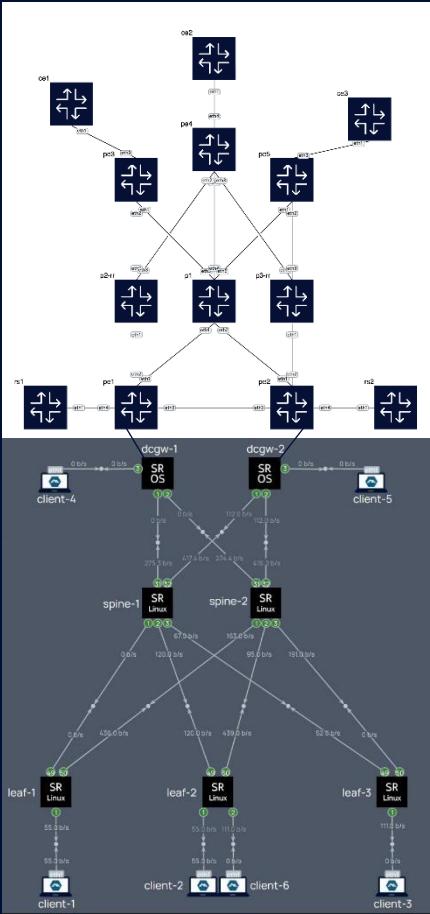
# What are we going to cover during this workshop?



1. Install containerlab
2. Container-based NOS lab
3. Startup configs
4. Working with Container registry
5. VM-based labs
6. Traffic capture
7. Getting our hands dirty and start some labs (...)

# What are we going to achieve by tomorrow end?

Gaining extensive knowledge on lab as code by building digital twins of:



1. A peering network similar to some of the largest IXPs in the world (SR-MPLS, EVPN architecture)
  2. An IXP private DC fabric for internal IT workloads using best design practices (IPv6 unnumbered, eBGP underlay, EVPN overlay, VXLAN)

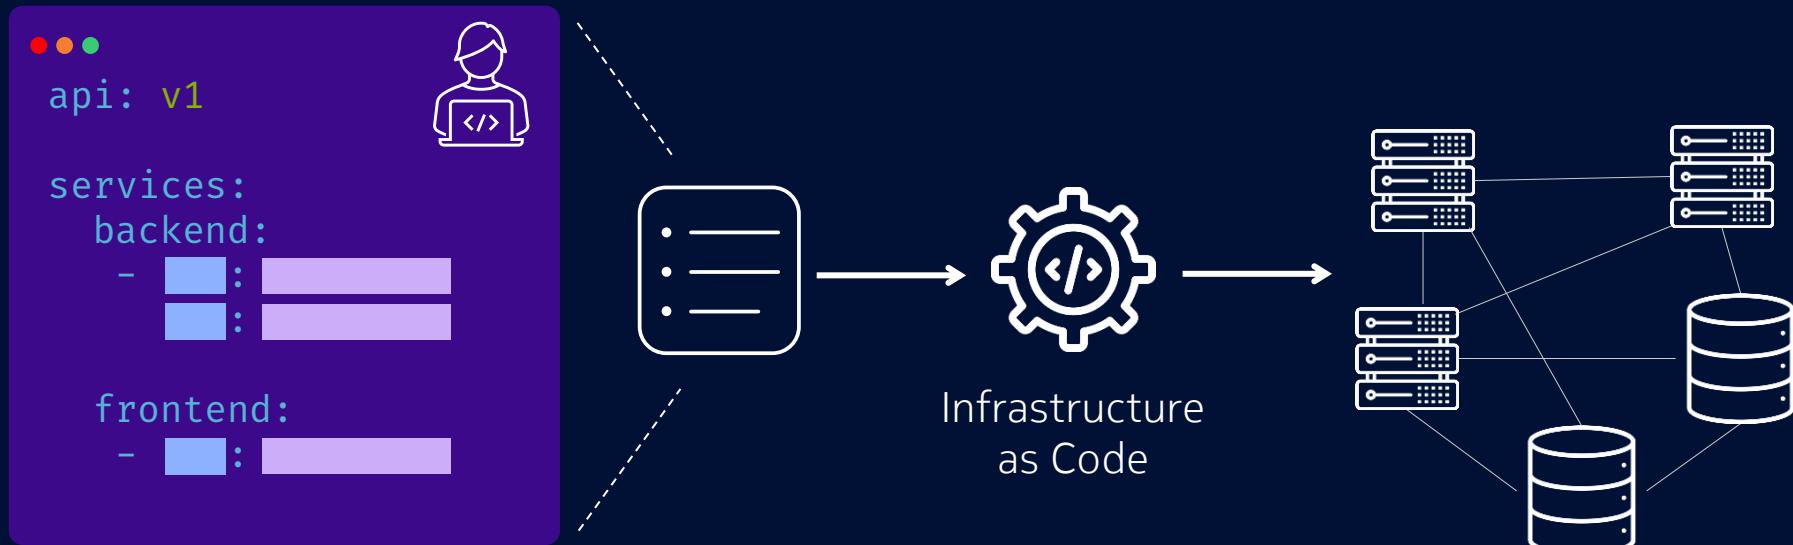
# Containerlab

“Lab as code” way to deploy networking labs



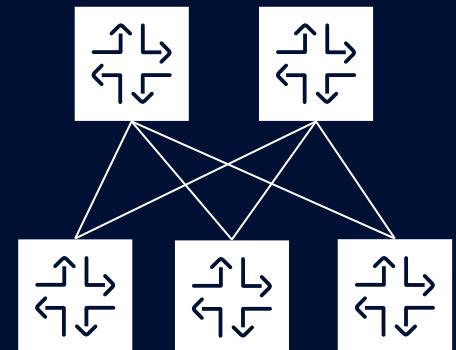
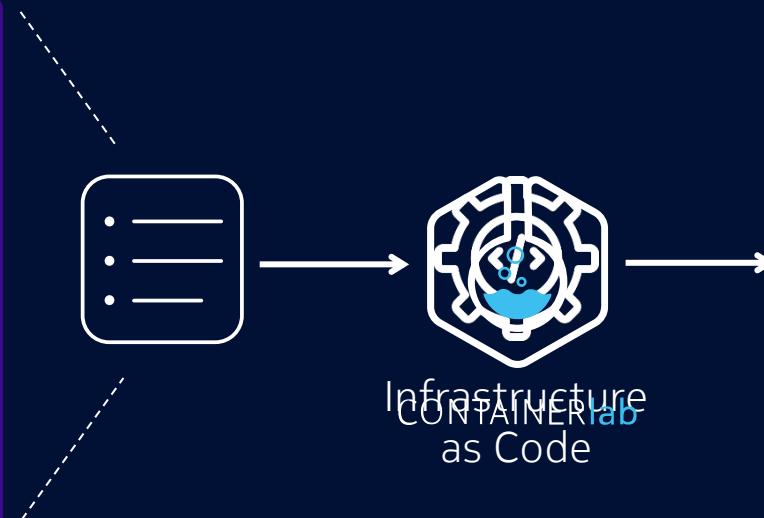
# How they deploy things over there?

Declarative, infrastructure as code approach



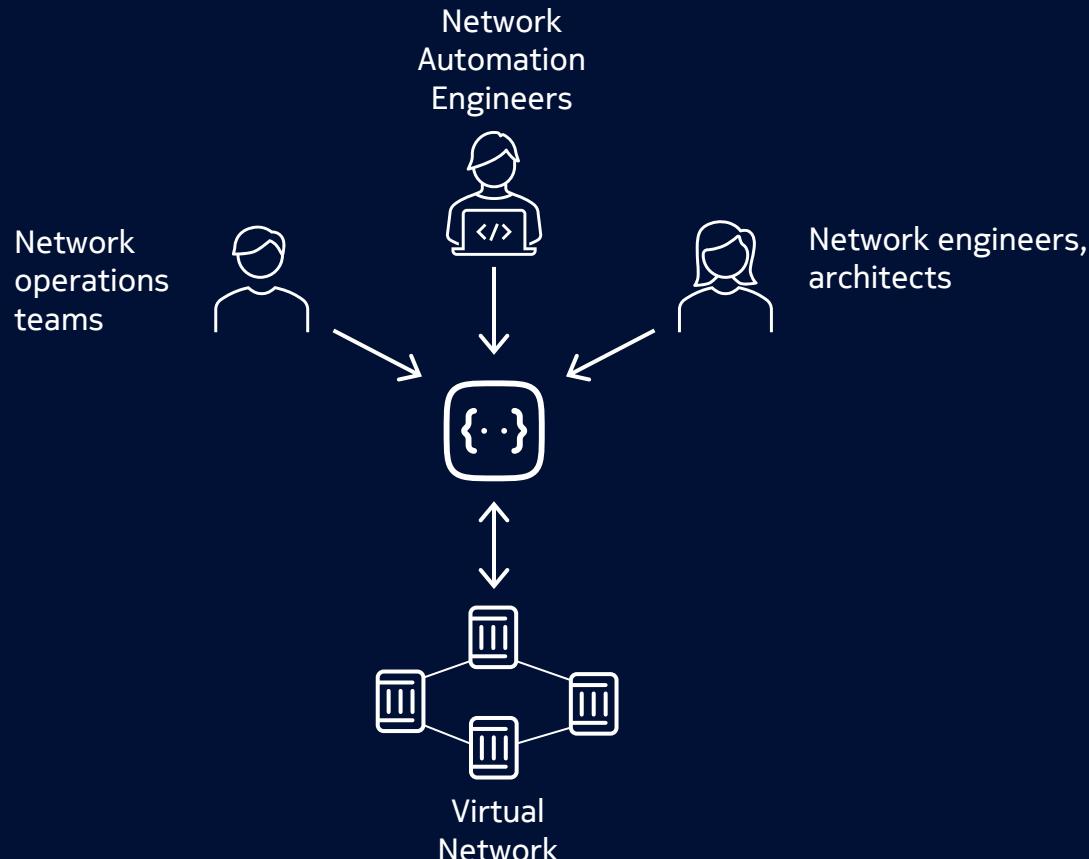
# Lab as code with Containerlab

```
...  
name: lab  
  
topology:  
  nodes:  
    - [ ]: [ ]  
    - [ ]: [ ]  
  
  links:  
    - [ ]: [ ]
```



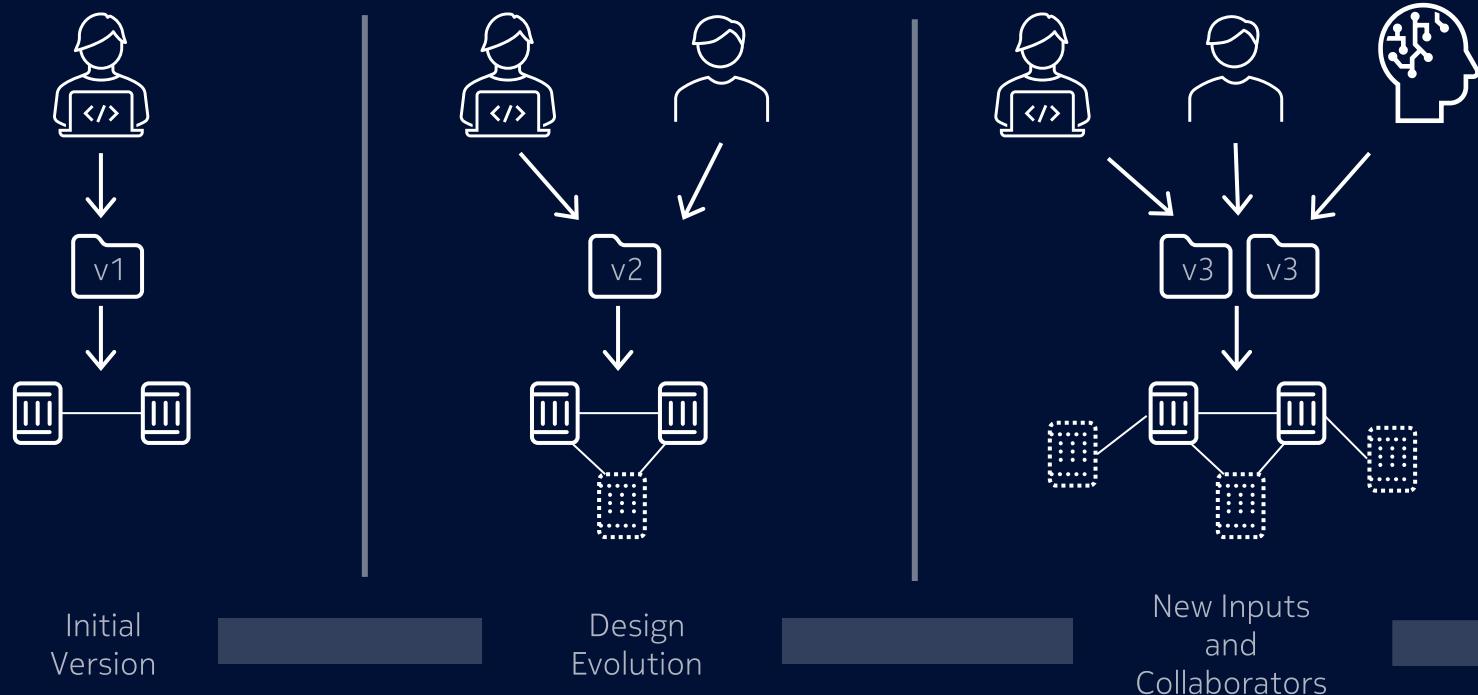
# Why Lab as Code?

## #1 - Collaboration



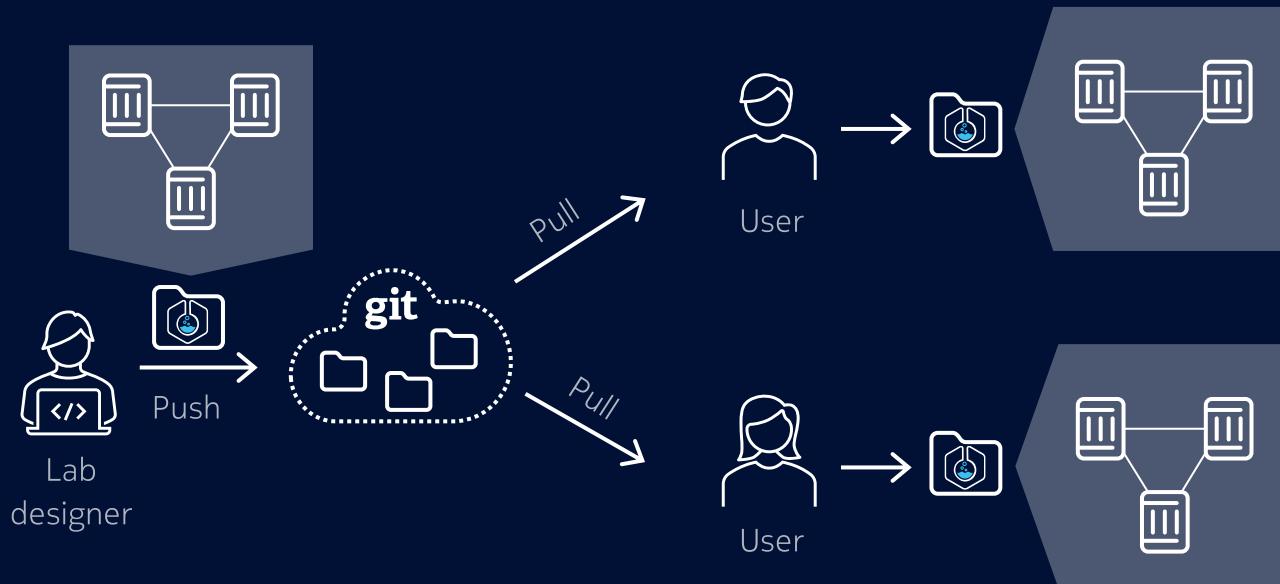
# Why Lab as Code?

## #2 - Versioning



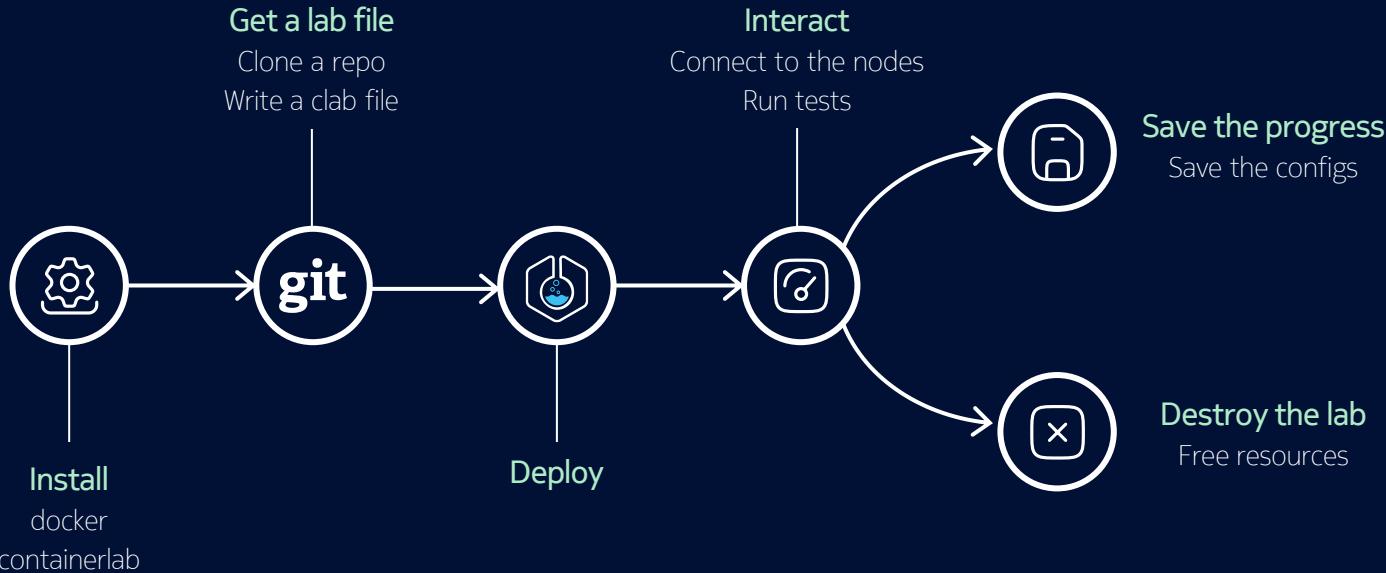
# Why Lab as Code?

## #3 - Sharing



# Containerlab

## The Workflow



# They use Containerlab



ARISTA



The New York Times



PACKETFABRIC

Proton

>>> network .toCode()

KEYSIGHT

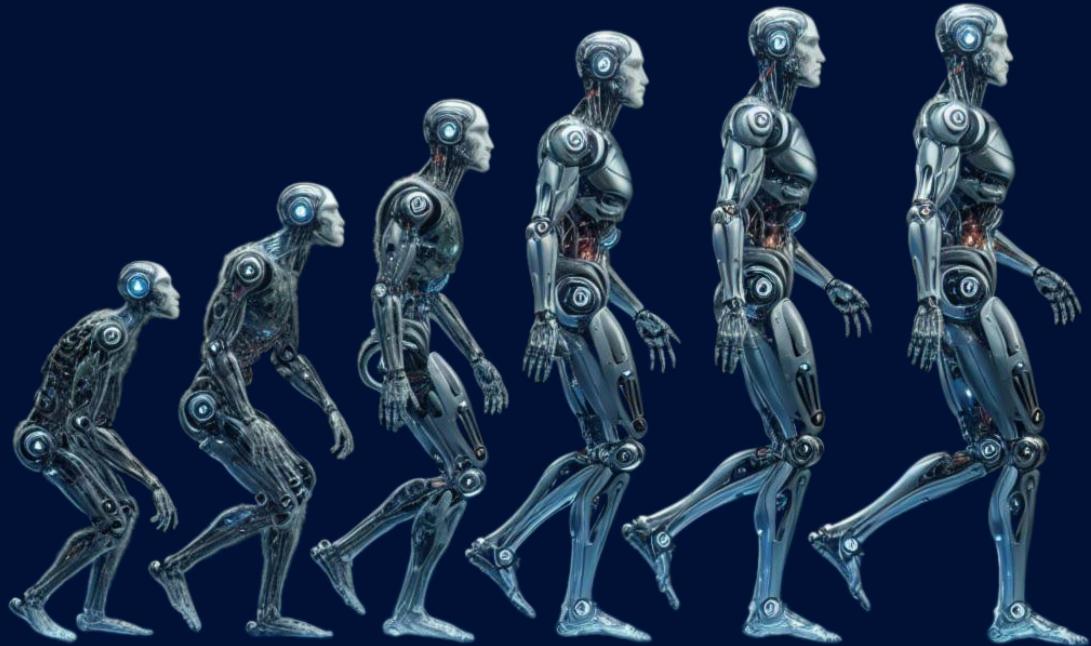


NOKIA

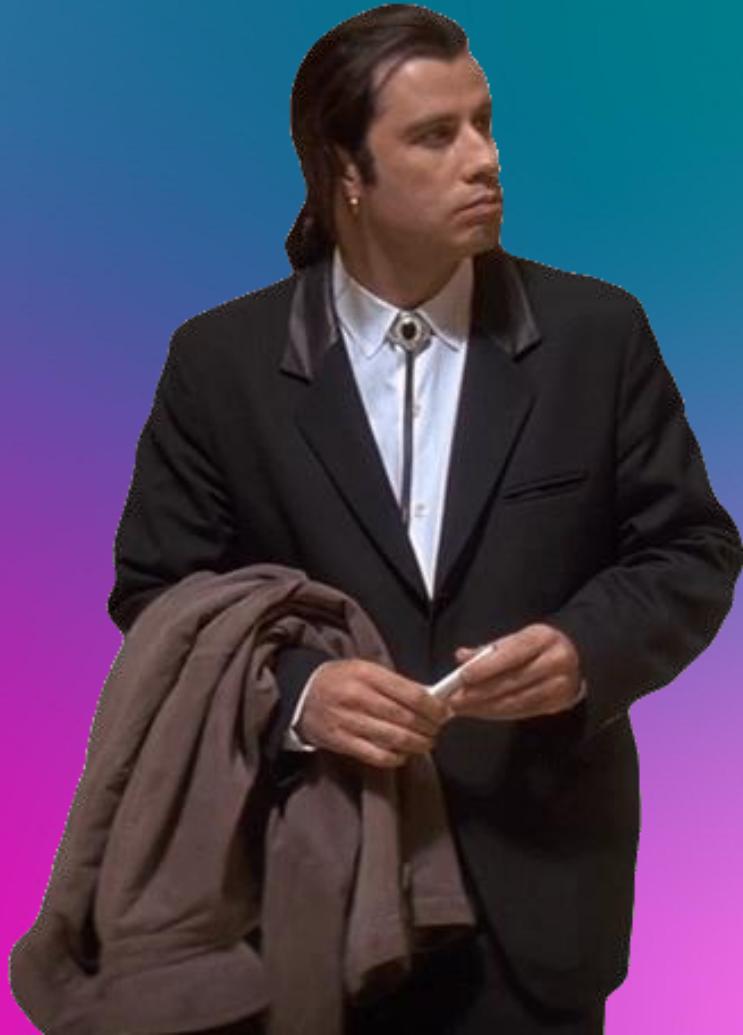
# Why not X?

**Changes in the environment  
drive the shifts in tooling.**

- Imperative -> Declarative
- Instructions -> Code
- Stateful -> Ephemeral
- One-off -> Repetitive



# Where is my Workshop?



NOKIA

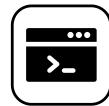
# Your personal hands-on experience

And how to get access to it



CONTAINERlab  
Workshop

Group ID:X



ssh: groupX@45.76.181.43

password: i8ClabW\$X



u: admin  
p: admin



u: admin  
p: admin



u: admin  
p: NokiaSrl1!



u: admin  
p: admin



[http://45.76.181.43:300X](http://45.76.181.43:300<u>X</u)

# Code of Conduct

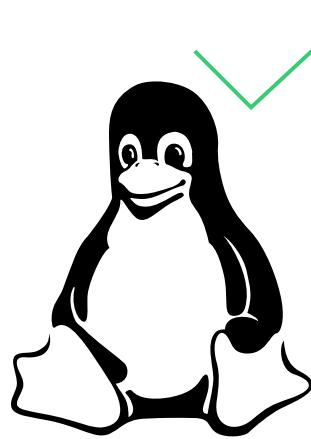
- Infra is precious,  
BEHAVE
- By looking at this slide  
you accept these simple  
rules. Easy.



# Installing containerlab



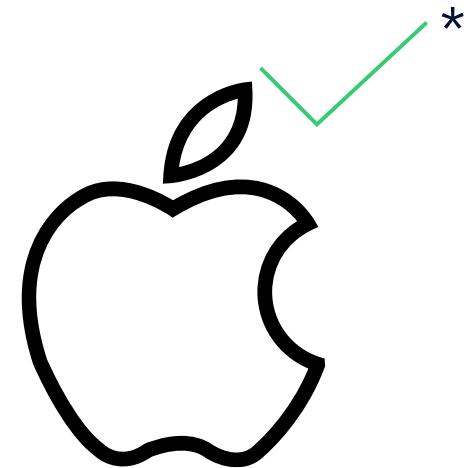
<https://containerlab.dev/install/>



LINUX



WSL2



MacOS

# All In One installer

The fastest way to get started

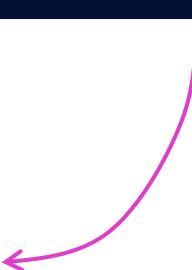


05-install

```
[*]-[<ID>]-[~]  
└─> curl -L https://containerlab.dev/setup | \  
    sudo bash -s "all"
```

## Installs:

- docker-ce
- containerlab
- GitHub CLI



**Log out and Log in**  
for group changes to take effect

# Installing containerlab



05-install

Verify install



```
[*]-[<ID>]-[~]
└─> docker run --rm hello-world
# Expected output: Hello from Docker!
```

```
[*]-[<ID>]-[~]
└─> containerlab version
```



```
version: 0.59.0
```

# Installing containerlab

## Other installation options



05-install

- Containerlab-only installation script
- APT/YUM/DNF package managers
- Binary via GitHub releases
- Container image with Containerlab

# Containerlab node types

Containerized Network OSes

- Sourced by the vendor
- Fast to spin up
- Small footprint
- Shareability and versioning

Current trend is to **move away**  
**from VM** packaging towards  
containers  
**for new NOses**

**NOKIA**  
SR Linux

**JUNIPER**  
NETWORKS  
cRPD

**ARISTA**  
cEOS

  
**CISCO**  
XRd

  
**NVIDIA**  
cVX

  
**KEYSIGHT**  
TECHNOLOGIES  
IXIA-c

and others...

# Containerlab node types

Regular container images

- All available container images
- Emulating clients
- Hundreds of network-focused software
  - Telemetry, logging stacks
  - Peering software
  - Flow collectors
  - BNG 
  - etc



Get / Set / Subscribe / Collect



Prometheus



NOKIA



10-basics

# Basic lab

# Basic lab

A simple starting point



## Topology definition

```
name: basic

topology:
  nodes:

    srl:
      kind: nokia_srlinux
      image: ghcr.io/nokia/srlinux:24.10.4

    ceos:
      kind: arista_ceos
      image: ceos:4.33.1F

    links:
      - endpoints: [srl:e1-1, ceos:eth1]
```

## Logical view



# Quickstart lab

Cloning the workshop repo



10-basics

```
cd ~  
git clone https://github.com/thcorre/innog8-workshop.git  
cd innog8-workshop/10-basic
```

Expected output:



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]  
└─> cat basic.clab.yml
```

# Containerlab Topology File

## Basic node definition

```
name: basic  
  
topology:  
  nodes:  
    srl:  
      kind: nokia_srlinux  
      image: ghcr.io/nokia/srlinux:24.10.4
```

1

Node definition container.  
Container name will be the node name.  
[Read more](#)

2

Kinds define the flavour of the node,  
it says if the node is a specific containerized  
Network OS or something else.  
[Read more](#)

3

Image specifies container image to use for  
this node.  
[Read more](#)



[topology definition file](#)

# Containerlab Topology File

## Links definition

### Topology definition

```
name: basic

topology:
  nodes:

    srl: <----+-----+
          |       |
          +-----+
    ceos: <----+-----+
          |       |
          +-----+
links:
  - endpoints: [srl:e1-1, ceos:eth1]
```

### Logical view



# Containerlab Topology File

Bringing nodes and links together

## Topology definition

```
name: basic

topology:
  nodes:

    srl:
      kind: nokia_srlinux
      image: ghcr.io/nokia/srlinux:24.10.4

    ceos:
      kind: arista_ceos
      image: ceos:4.33.1F

  links:
    - endpoints: [srl:e1-1, ceos:eth1]
```

## Logical view



# Quickstart

## Deploying the lab



10-basics



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]  
└─> sudo containerlab deploy -t basic.clab.yml
```



Containerlab pulls container images specified in the topology file



Error response from daemon: pull access denied for **ceos**, repository does not exist or may require 'docker login': denied: requested access to the resource is denied

# Local container image store



10-basics

## Topology definition

```
srl:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux:24.10.4
```

```
ceos:  
  kind: arista_ceos  
  image: ceos:4.33.1F
```

## Local image store



[\*]—[<ID>]—[~/innog8-workshop/10-basic]

└─> **docker images**

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hello-world	latest	d2c94e258dcb	11 months ago	13.3kB

# Container image notation

```
image: ghcr.io/nokia/srlinux
```

```
image: ceos:4.33.1F
```

```
image: prom/prometheus:v2.47
```



## Fully qualified name:

ghcr.io - registry name  
nokia - org name  
srlinux - repo name  
latest - implicit repo tag

## Fully qualified name:

docker.io - implied registry name  
library - implied org name  
ceos - repo name  
4.33.1F - repo tag

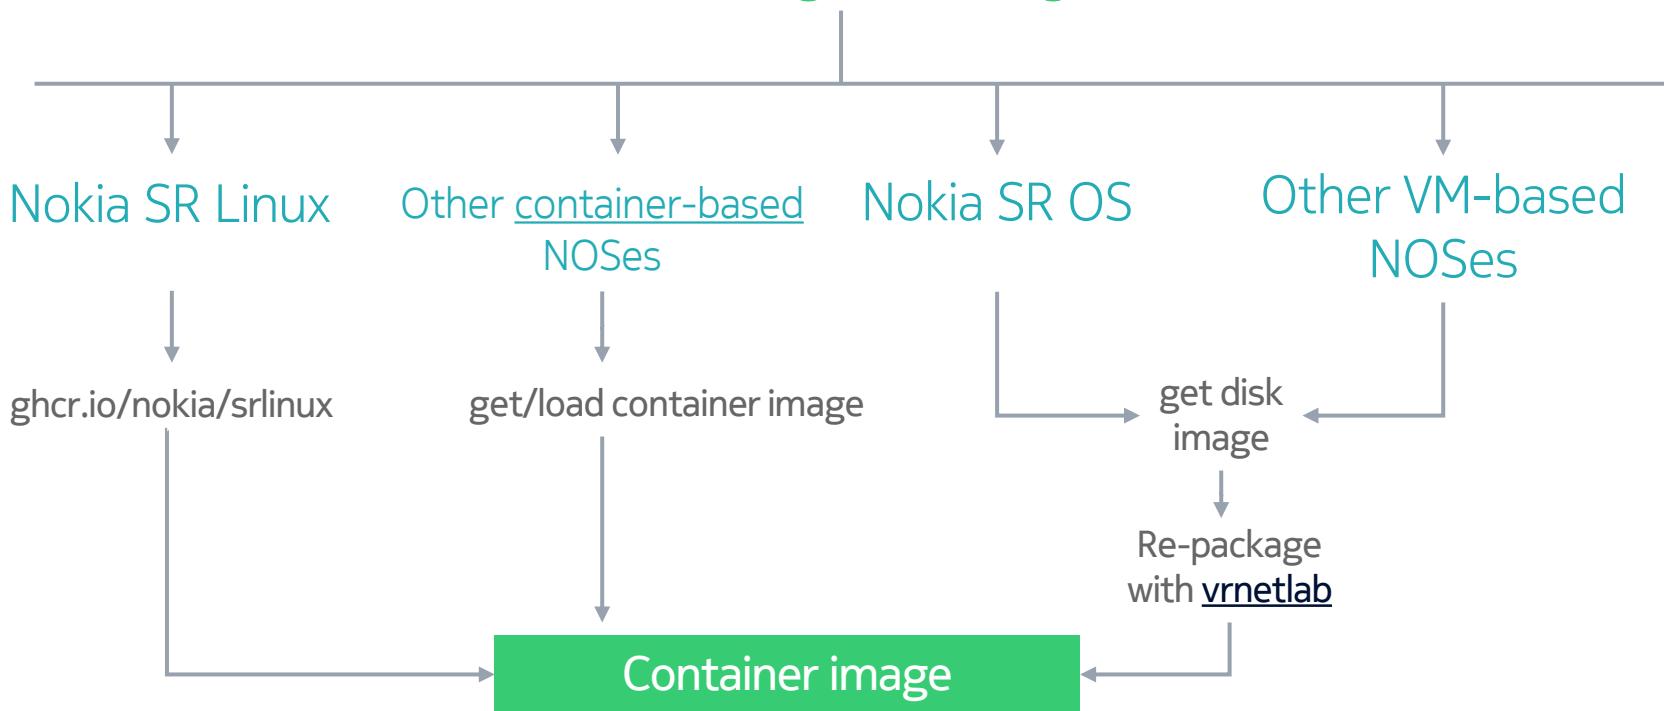
## Fully qualified name:

docker.io - implied registry name  
prom - implied org name  
prometheus - repo name  
v2.47 - repo tag

# Container images

Where do I get one?

## How do I get an image?



# Pulling the public image



10-basics

SR Linux container image



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]
└─> docker pull ghcr.io/nokia/srlinux
Using default tag: latest
latest: Pulling from nokia/srlinux
1411e7dd712e: Downloading [=====] 312.7MB/874.8MB
```



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]
└─> docker images
REPOSITORY          TAG      IMAGE ID   CREATED    SIZE
ghcr.io/nokia/srlinux  latest   36427a3c297e  3 weeks ago  3.03GB
hello-world         latest   d2c94e258dcb  11 months ago 13.3kB
```

# Importing Arista cEOS image



[containerlab.dev/manual/kinds/ceos/](https://containerlab.dev/manual/kinds/ceos/)



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]
└─> docker import ~/images/cEOS64-lab-4.33.1F.tar.xz ceos:4.33.1F
--no output for some time-
```

```
sha256:1ea1f7ed3695e0d9abf356d6a530293f319fcc148716c93503b5eabaaf087aa
```

# Checking local image store



10-basics

After pulling and importing container images

## Topology definition

```
srl:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux  
  
ceos:  
  kind: arista_ceos  
  image: ceos:4.33.1F
```

## Local image store



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]  
└── docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
ceos	4.33.1F	abc71bccbc4f	5 seconds ago
ghcr.io/nokia/srlinux	latest	36427a3c297e	3 weeks ago

# Quickstart lab

Retry deploying the lab



10-basics



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]  
└─> sudo containerlab deploy -t basic.clab.yml
```

```
INFO[0031] Adding ssh config for containerlab nodes
```

#	Name	Container ID	Image	Kind	State	IPv4 Address	IPv6 Address
1	clab-basic-ceos	9d0b5121a2ff	ceos:4.33.1F	arista_ceos	running	172.20.20.3/24	2001:172:20:20::3/64
2	clab-basic-srl	279d170737d6	ghcr.io/nokia/srlinux	nokia_srlinux	running	172.20.20.2/24	2001:172:20:20::2/64

# Quickstart lab



## Connecting to the nodes

#	Name	Kind	IPv4 Address	IPv6 Address
1	clab-basic-ceos	arista_ceos	172.20.20.3/24	2001:172:20:20::3/64
2	clab-basic-srl	nokia_srlinux	172.20.20.2/24	2001:172:20:20::2/64

A dark blue terminal window with three colored dots (red, orange, green) at the top. Below them, the text "ssh clab-basic-srl" is displayed in white.A dark blue terminal window with three colored dots (red, orange, green) at the top. Below them, the text "ssh admin@172.20.20.3" is displayed in white.

# Containerlab /etc/hosts automation



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]
└─> cat /etc/hosts
127.0.0.1    localhost
::1          localhost ip6-localhost ip6-loopback

##### CLAB-basic-START #####
172.20.20.3      clab-quick-ceos
172.20.20.2      clab-quick-srl
2001:172:20:20::3    clab-quick-ceos
2001:172:20:20::2    clab-quick-srl
##### CLAB-basic-END #####
```

# Containerlab ssh config automation



```
[*]-[<ID>]-[~/innog8-workshop/10-basics]
└─> cat /etc/ssh/ssh_config.d/clab-basic.conf
# Containerlab SSH Config for the quick lab
Host clab-basic-ceos
User admin
StrictHostKeyChecking=no
UserKnownHostsFile=/dev/null

Host clab-basic-srl
User admin
StrictHostKeyChecking=no
UserKnownHostsFile=/dev/null
```

# Default node configuration

- Management interfaces enabled  
gNMI, Netconf, REST API, SNMP, etc
- TLS certificates created\*
- LLDP enabled\*
- SSH Keys\*
- Documented in the kind documentation

\*for some nodes

## Node configuration #

SR Linux uses a `/etc/opt/srlinux/config.json` file to persist its configuration. By default, containerlab starts nodes of `srl` kind with a basic "default" config, and with the `startup-config` parameter, it is possible to provide a custom config file that will be used as a startup one.

### Default node configuration

When a node is defined without the `startup-config` statement present, containerlab will make **additional configurations** on top of the factory config:

```
# example of a topo file that does not define a custom startup-config
# as a result, the default configuration will be used by this node
```

```
name: srl_lab
topology:
  nodes:
    srl1:
      kind: nokia_srlinux
      type: ixrd3
```



[containerlab.dev/manual/kinds/srl/#node-configuration](https://containerlab.dev/manual/kinds/srl/#node-configuration)

# Checking the network connectivity



ssh clab-basic-srl



ssh clab-basic-ceos

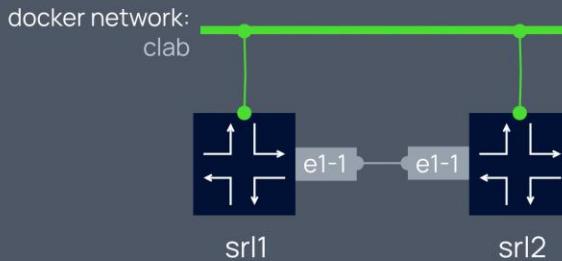
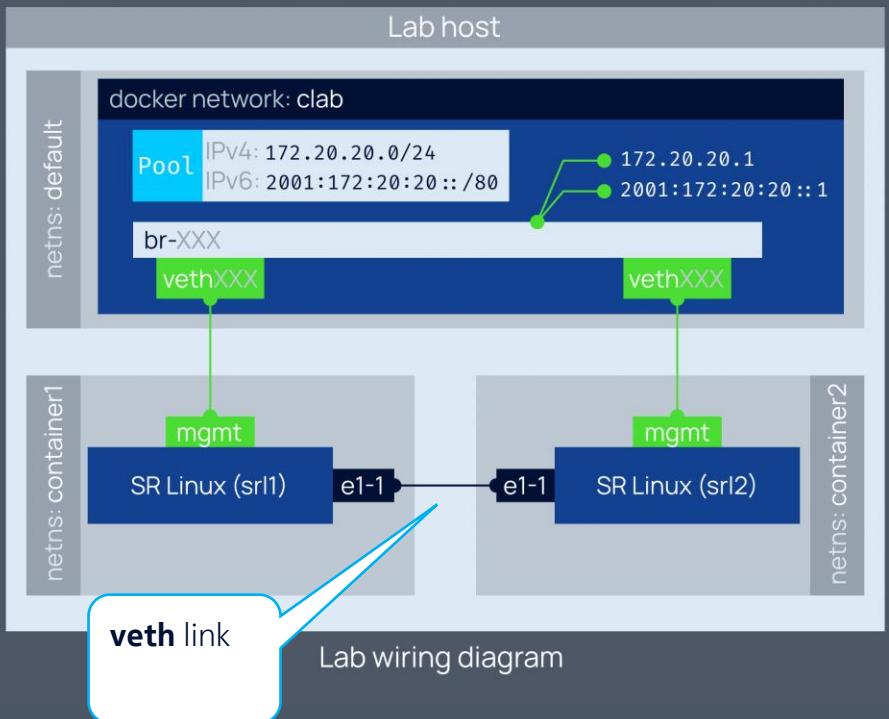


show /system lldp neighbor interface ethernet-1/1

Name	Neighbor	Neighbo r System Name	Neighbo r Chassis ID	Neighbo r First Message ago	Neighbo r Last Update ago	Neighbo r Port t1
ethernet -1/1	00:1C:73 :46:95:5 C	ceos	00:1C:7 3:46:95 :5C	19 hours ago	a second ago	Etherne t1



# Containerlab networking



# Listing running labs

Command	Effect
<code>sudo containerlab inspect [--topo &lt;path to clab file&gt;]</code>	List nodes of a lab found in the \${PWD} or by the --topo path.
<code>sudo clab ins [-t &lt;path&gt;]</code>	
<code>sudo containerlab inspect --all</code>	List all deployed labs and their nodes
<code>sudo clab ins -a</code>	



```
sudo clab inspect -a
```

# Lab directory

a.k.a Persistent Storage

```
> tree -L 3 clab-basic
clab-demo1
├── ansible-inventory.yml
├── authorized_keys
└── ceos
    └── flash
        ├── startup-config
        └── boot-config
└── srl
    └── config
        └── config.json
└── topology-data.json
```



[configuration artifacts](#)

Lab directory name: **clab-basic**

fixed prefix  
lab name

## Lab directory:

- Takes precedence over the startup-config
- Better not be checked into git
- Use startup-config instead of relying on the config in the lab dir
- Use startup-config instead of relying on the lab-directory

# Destroying the lab



10-basics

Command	Effect
<code>sudo containerlab <b>destroy</b> [--topo &lt;path to clab file&gt;]</code>	Removes containers for a given topology. <b>Leaves</b> a lab directory intact
<code>sudo containerlab <b>destroy</b> [--topo &lt;path to clab file&gt;] --cleanup</code>	Removes containers <b>and a lab directory</b> for a given topology.
<code>sudo containerlab <b>destroy</b> --all</code>	Remove containers for all running labs. Keep lab directories.
<code>sudo containerlab <b>destroy</b> --all --cleanup</code>	Remove containers and lab directories for all running labs.



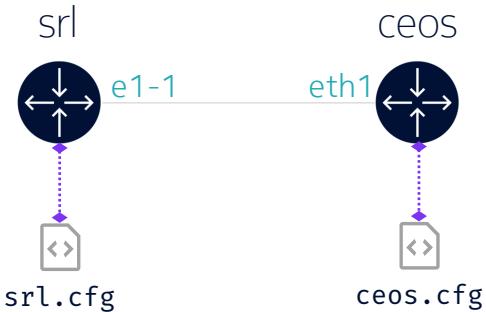
[\*]-[<ID>]-[~/innog8-workshop/10-basics]  
└→ **sudo clab des --cleanup**

# Startup configuration

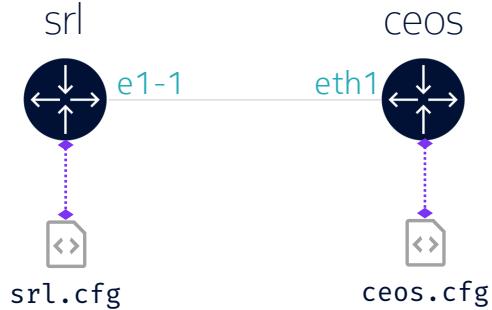
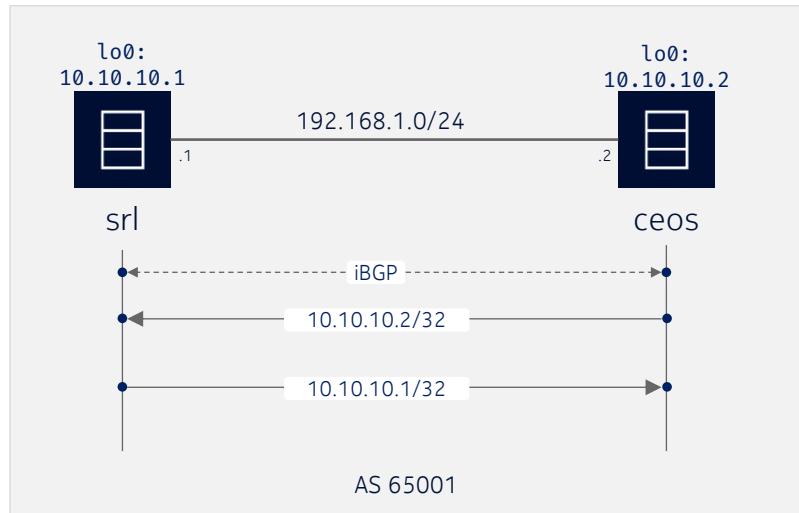
# Startup configuration

## Startup configuration

- The best way to provide the initial config
- Should be checked into git to keep the lab fully declarative
- Typically contained in an external file
- Provided in the CLI syntax and/or the serialization format of the configuration
- [Check the Kind documentation for details](#)



# Startup configuration



# Providing startup configuration

## topology definition

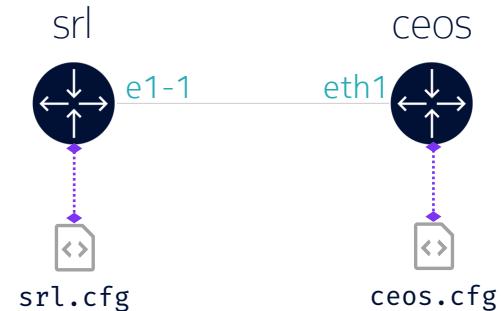
```
name: startup                                         startup.clab.yml

topology:
  nodes:

    srl:
      kind: nokia_srlinux
      image: ghcr.io/nokia/srlinux
      startup-config: srl.cfg

    ceos:
      kind: arista_ceos
      image: ceos:4.33.1F
      startup-config: ceos.cfg
```

## logical view



# What is inside the startup config?

srl.cfg

```
interface ethernet-1/1 {  
    subinterface 0 {  
        admin-state enable  
        ipv4 {  
            admin-state enable  
            address 192.168.1.1/24 {  
            }  
        }  
    }  
}  
  
interface lo0 {  
    subinterface 0 {  
        admin-state enable  
        ipv4 {  
            address 10.10.10.1/32 {  
            }  
        }  
    }  
}
```

ceos.cfg

```
management api http-commands  
    no shutdown  
{{- if .Env.CLAB_MGMT_VRF }}  
!  
    vrf {{ .Env.CLAB_MGMT_VRF }}  
        no shutdown  
{{end}}  
!  
  
interface Ethernet1  
    no switchport  
    ip address 192.168.1.2/24  
!  
interface Loopback0  
    ip address 10.10.10.2/32  
!
```

# Startup config lab



15-startup



```
0 [*]-[<ID>]-[~/innog8-workshop/10-basics]  
└─> cd ~/innog8-workshop/15-startup
```

# Deploy and see the effect of the startup config



1



```
[*]-[<ID>]-[~/innog8-workshop/15-startup]  
└─> sudo clab dep -c
```

**-c | --reconfigure**  
(destroys the lab before deploy)

2



```
[*]-[<ID>]-[~/innog8-workshop/15-startup]  
└─> ssh clab-startup-srl
```

3



```
--{ running }--[ ]--  
A:srl# show network-instance default protocols bgp neighbor 192.168.1.2
```

# Deploy and see the effect of the startup config

5

```
● ● ●  
[*]-[<ID>]-[~/innog8-workshop/15-startup]  
└─> ssh clab-startup-ceos
```

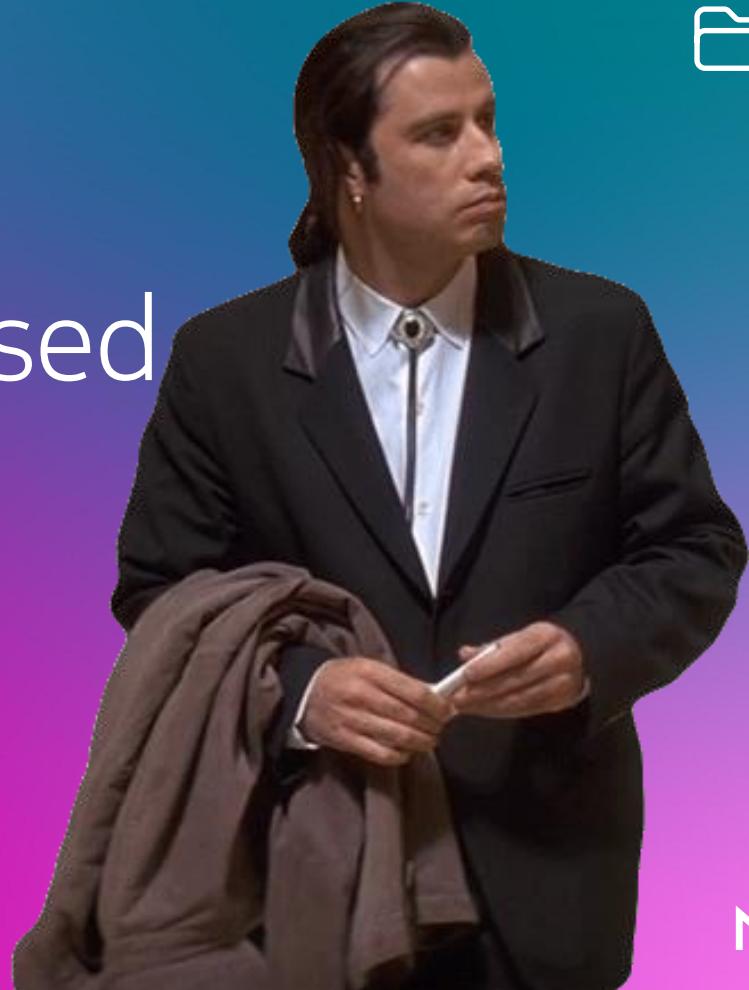
6

```
● ● ●  
ceos>ping 10.10.10.1
```

7

```
● ● ●  
[*]-[<ID>]-[~/innog8-workshop/15-startup]  
└─> docker exec -t clab-startup-ceos Cli -c 'ping 10.10.10.1'
```

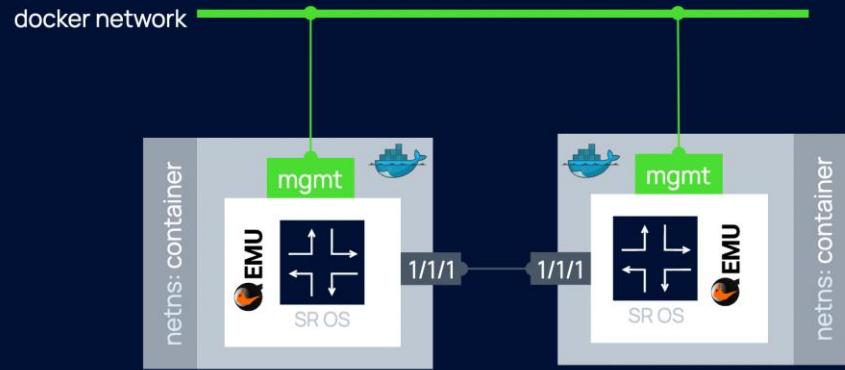
Where are my  
Good old VM-based  
Network OSes?



# Containerlab node types

Virtual machines in a container package

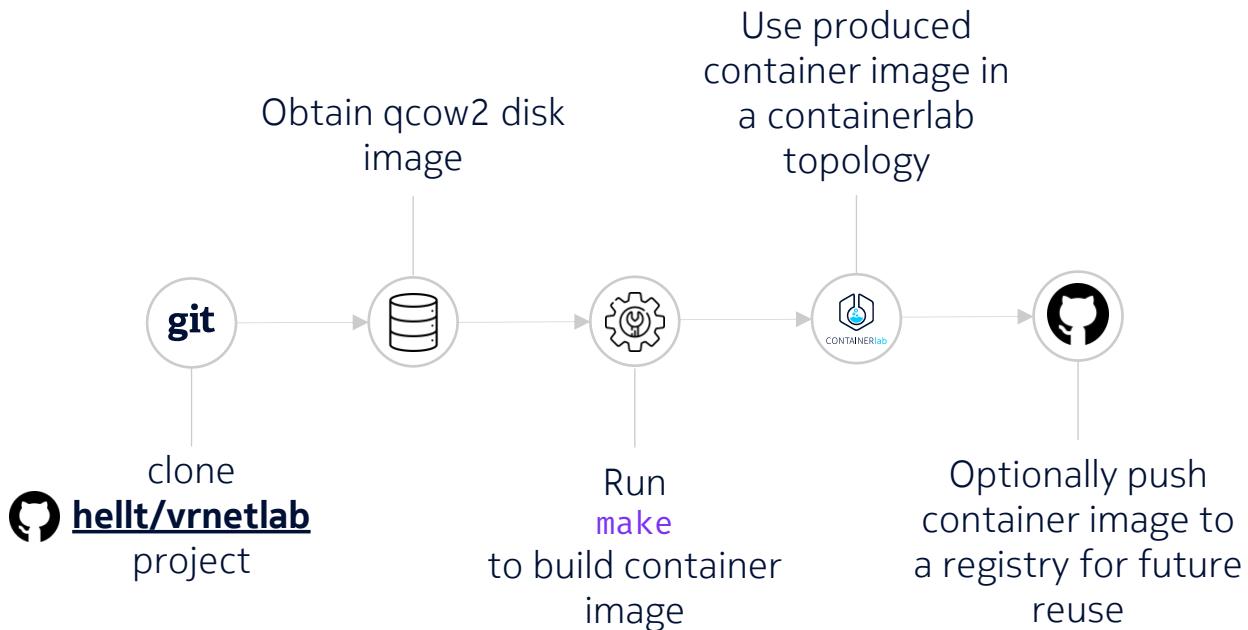
- Traditional Network OS packaged as a VM
- Integrated with containerlab via [vrnetlab](#) open-source project
- Onboard existing VM-based NOSes



Juniper vMX, vSwitch, EVO  
Arista vEOS  
Cisco XRv9k, c8000v, NX-OS  
Fortinet Fortigate  
IPInfusion OcNOS  
Palo Alto PAN-OS  
Huawei VRP,  
And more!

# Bringing VM-based nodes to Containerlab

## The workflow



# VM-based nodes topology

## topology definition

```
name: vm  
  
topology:  
  nodes:  
  
    sonic:  
      kind: sonic-vm  
      image: vrnetlab/sonic-vm:202411  
  
    srl:  
      kind: nokia_srlinux  
      image: ghcr.io/nokia/srlinux:24.10.4  
  
  links:  
    - endpoints: [sonic:eth1, srl:e1-1]
```

## logical view



NEW

# Interface aliases

## topology definition

```
links:  
  - endpoints: [sonic:eth1, srl:e1-1]
```

||

```
links:  
  - endpoints: [sonic:eth1, srl:eth1]
```

## logical view



# Cloning hellt/vrnetlab



20-vm

```
● ● ●
[*]-[<ID>]-[~/innog8-workshop/15-startup]
└──> cd ~
git clone https://github.com/hellt/vrnetlab.git
cd ~/vrnetlab
```

# Building Sonic image 1/2



```
● ● ●  
[*]-[<ID>]-[~/vrnetlab]  
└─> cp ~/images/sonic-vs.img.qcow2  
~/vrnetlab/sonic-vs/
```

```
● ● ●  
[*]-[<ID>]-[~/vrnetlab]  
└─> cd ~/vrnetlab/sonic-vs  
make
```

# Building Sonic image 2/2



```
[*]-[<ID>]-[~/vrnetlab]
  └── docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
vrnetlab/sonic-vm	202411	6bb8d92e8a12	10 seconds ago	6.35GB
ceos	4.33.1F	42f52d70fc57	40 minutes ago	2.46GB
ghcr.io/nokia/srlinux	24.10.4	eb2a823cd8ce	8 days ago	2.14GB

# Building SR OS container image 1/2



```
● ● ●  
[*]-[<ID>]-[~/vrnetlab]  
└─> cp ~/images/sros-vm-24.10.R4.qcow2 ~/vrnetlab/sros/
```

# Building SR OS container image 2/2



```
● ● ●
[*]-[<ID>]-[~/vrnetlab]
└─> cd ~/vrnetlab/sros
     make

=> => naming to docker.io/vrnetlab/nokia_sros:24.10.R4
```

```
● ● ●
[*]-[<ID>]-[~/vrnetlab]
└─> docker images | grep sros

vrfnetlab/nokia_sros    24.10.R4    0a9146ccdd14    19 minutes ago    1.43GB
```

# Deploying the lab



```
[*]-[<ID>]-[~]  
└─> cd ~/innog8-workshop/20-vm  
      sudo clab dep -c
```

# Monitoring the boot process

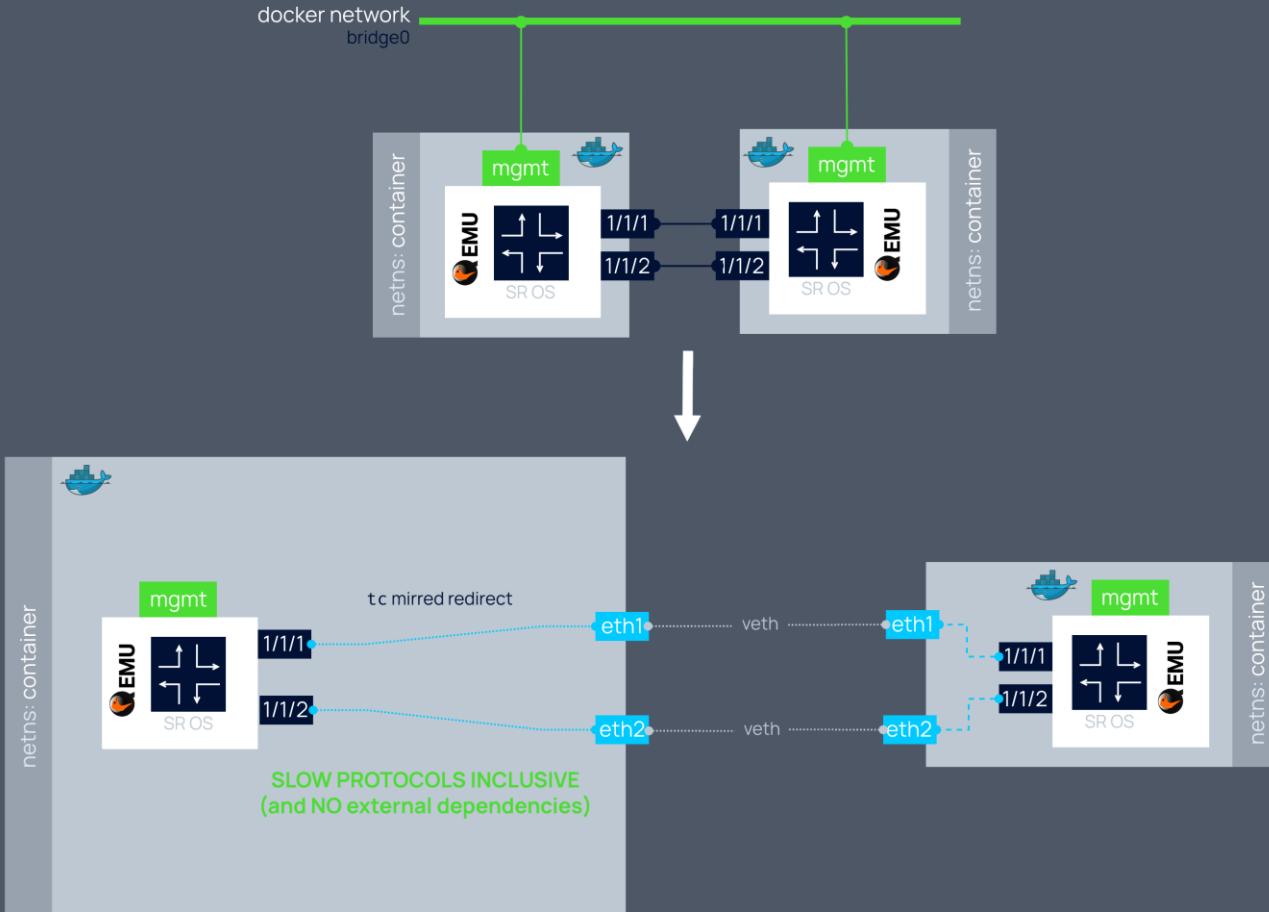


```
● ● ●  
[*]-[<ID>]-[~]  
└─> docker logs -f clab-vm-sonic
```

## Boot log

- The first thing to check when something “does not work”
- Gives insight into what is provisioned by Containerlab in terms of the base configuration

# VM nodes under the hood



# Connecting to the nodes



Name	Kind/Image	State	IPv4/IPv6 Address
clab-vm-sonic	sonic-vm vrnetlab/sonic-vm:202411	running (healthy)	172.20.20.2 2001:172:20:20::2
clab-vm-srl	nokia_srlinux ghcr.io/nokia/srlinux:24.10.4	running	172.20.20.3 2001:172:20:20::3



A dark blue rectangular box representing a terminal window. It contains three colored dots (red, orange, green) at the top left and the command "ssh clab-vm-srl" in white text below them.

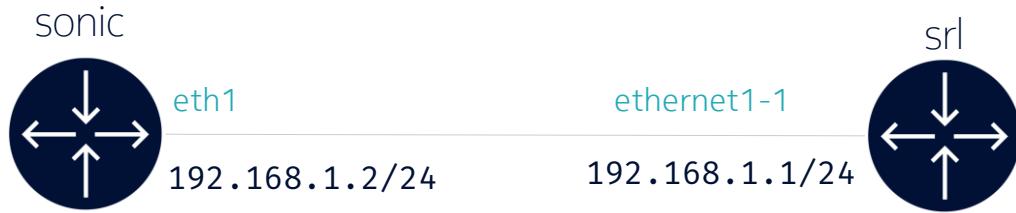


A dark blue rectangular box representing a terminal window. It contains three colored dots (red, orange, green) at the top left and the command "ssh clab-vm-sonic" in white text below them.

# Configure and verify the basic L3 setup



Refer to the workshop repo

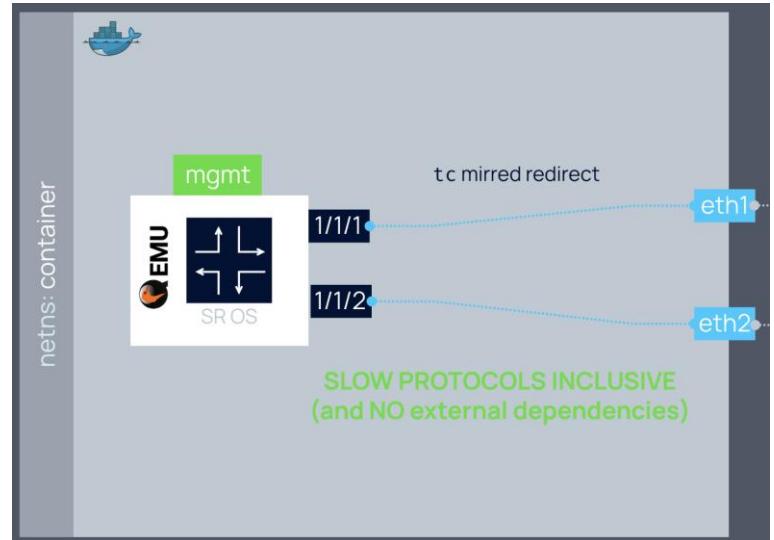


- Verify L3 config
- Verify datapath

# Networking concepts of the VM lab nodes



- Container interfaces (eth0, eth1, etc) are **not** the interfaces of the embedded VM.
- **tapX** interfaces are the VM's interfaces
- Data interfaces use **tc** to stitch container interfaces with VM interfaces
- Management interfaces use qemu user networking



# Networking concepts of the VM lab nodes

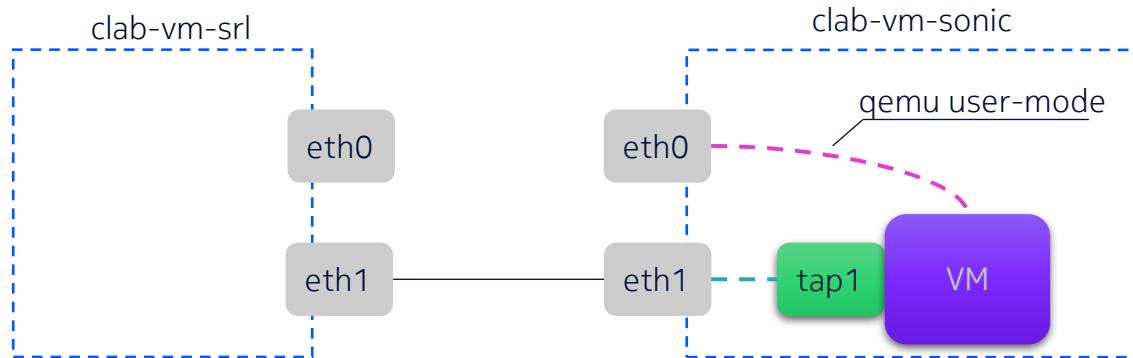
20-vm

```
[*]-[<ID>]-[~]  
└─> docker exec -it clab-vm-sonic bash
```

- Enter container shell

```
root@sonic:/# ip link
```

- List Linux interfaces

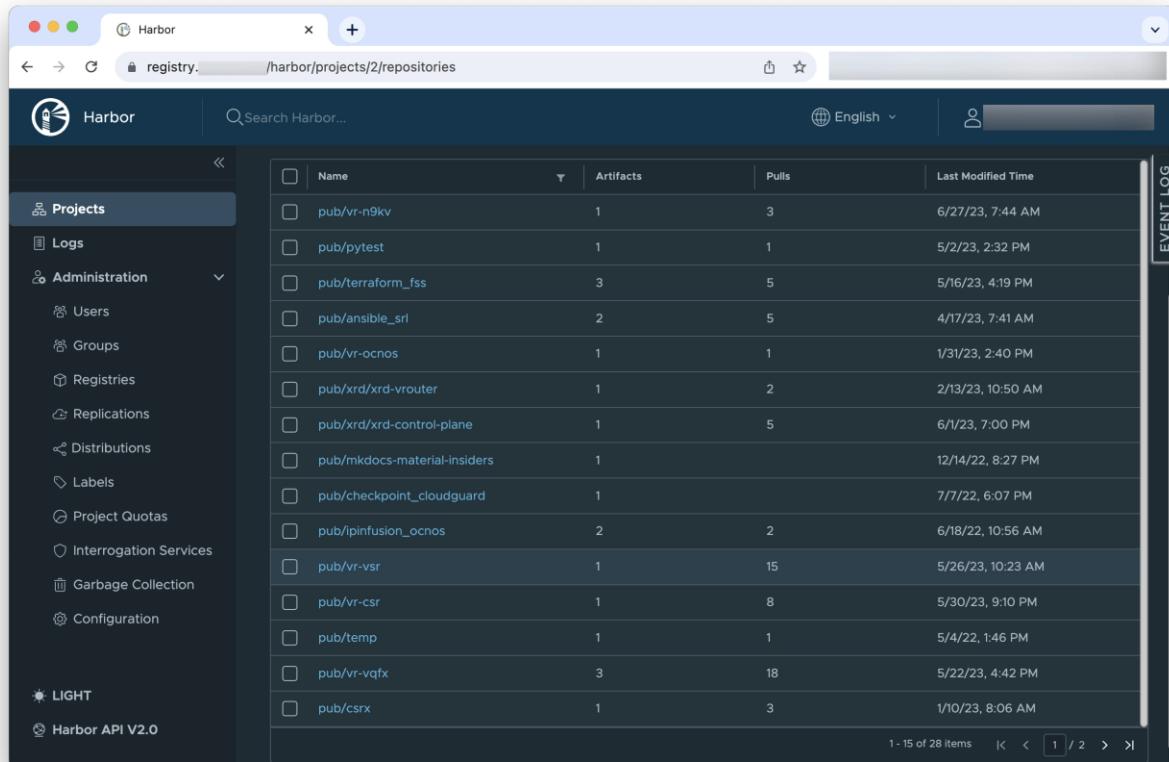


# Container registry

# Container registry

Taking back control over your images

- Centralized image repository
- Version control (tags, SHA)
- Seamless containerlab integration
- Granular access control



The screenshot shows the Harbor Container Registry web interface. The left sidebar contains navigation links: Projects, Logs, Administration (with sub-links for Users, Groups, Registries, Replications, Distributions, Labels, Project Quotas, Interrogation Services, Garbage Collection, and Configuration), LIGHT mode, and Harbor API V2.0. The main content area displays a table of repositories. The table has columns for Name, Artifacts, Pulls, and Last Modified Time. There are 28 items listed, with the first few being pub/vr-n9kv, pub/pytest, pub/terraform\_fss, pub/ansible\_sri, pub/vr-ocnos, pub/xrd/xrd-vrouter, pub/xrd/xrd-control-plane, pub/mkdocs-material-insiders, pub/checkpoint\_clouguard, pub/ipinfusion\_ocnos, pub/vr-vsr, pub/vr-csr, pub/temp, pub/vr-vqfx, and pub/csrx.

Name	Artifacts	Pulls	Last Modified Time
pub/vr-n9kv	1	3	6/27/23, 7:44 AM
pub/pytest	1	1	5/2/23, 2:32 PM
pub/terraform_fss	3	5	5/16/23, 4:19 PM
pub/ansible_sri	2	5	4/17/23, 7:41 AM
pub/vr-ocnos	1	1	1/31/23, 2:40 PM
pub/xrd/xrd-vrouter	1	2	2/13/23, 10:50 AM
pub/xrd/xrd-control-plane	1	5	6/1/23, 7:00 PM
pub/mkdocs-material-insiders	1		12/14/22, 8:27 PM
pub/checkpoint_clouguard	1		7/7/22, 6:07 PM
pub/ipinfusion_ocnos	2	2	6/18/22, 10:56 AM
pub/vr-vsr	1	15	5/26/23, 10:23 AM
pub/vr-csr	1	8	5/30/23, 9:10 PM
pub/temp	1	1	5/4/22, 1:46 PM
pub/vr-vqfx	3	18	5/22/23, 4:42 PM
pub/csrx	1	3	1/10/23, 8:06 AM

# Harbor container registry

## Web access



<https://45.76.181.43/>

admin: i8ClabW\$

The screenshot shows the Harbor web interface. On the left is a sidebar with options like Projects, Logs, Administration, Users, Robot Accounts, Registries, Replications, Distributions, Labels, Project Quotas, Interrogation Services, Clean Up, Job Service Dashboard, and Configuration. The main area has a header with 'Search Harbor...', language ('English'), and user ('admin'). It displays summary statistics: 0 Private and 1 Public projects, 0 Private and 2 Public repositories, and a total quota used of 1.64 GiB. Below this is a table of projects, showing one entry: 'library' (Public, Project Admin, Project, 2 repositories, created 4/11/24, 5:55 PM). The table includes columns for Project Name, Access Level, Role, Type, Repositories Count, and Creation Time. A search bar and a page size dropdown (15) are also present.



<https://45.76.181.43/>

# Listing available images

Harbor UI

The screenshot shows the Harbor UI interface. The top navigation bar includes the Harbor logo, a search bar with the placeholder "Search Harbor...", language and region dropdowns (English, Default), and a user account for "admin". The left sidebar has a "Projects" section selected, along with other options like "Logs", "Administration", "Users", "Robot Accounts", "Registries", "Replications", "Distributions", "Labels", "Project Quotas", and "Interrogation Services". The main content area displays a summary for a project named "lib...". It shows the "Access Level" as "Public" and the "Quota used" as "670.69MiB of unlimited". Below this, tabs for "Summary", "Repositories", "Members", "Labels", "Scanner", and "..." are present, with "Repositories" being the active tab. A "PUSH COMMAND" button and a search bar are also visible. A table lists the available images, with one entry highlighted: "library/nokia\_sros" (1 artifact, 0 pulls, last modified 11/7/24, 3:57 PM). The bottom of the screen shows pagination controls for "Page size" (15) and "1 - 1 of 1 items".

Name	Artifacts	Pulls	Last Modified Time
library/nokia_sros	1	0	11/7/24, 3:57 PM

# Listing available images

With API

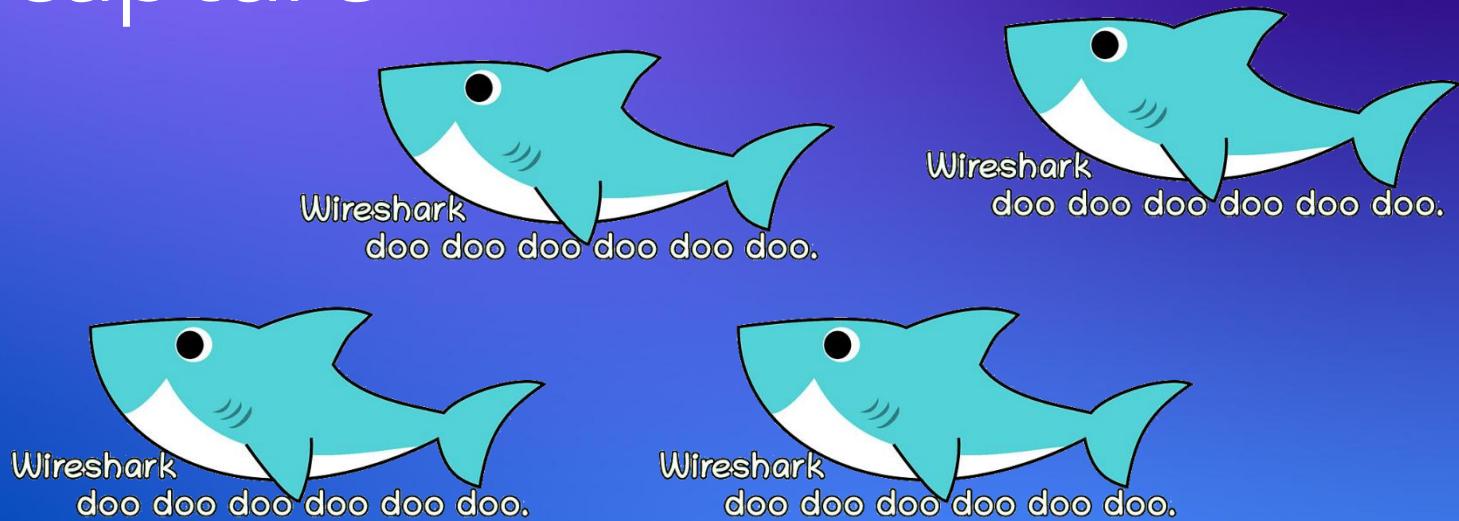


30-registry

Listing repositories

```
● ● ●
[*]-[<ID>]-[~]
└─> curl -X 'GET'
'https://45.76.181.43/api/v2.0/repositories?page=1&page_size=10' -H 'accept:
application/json' -k
```

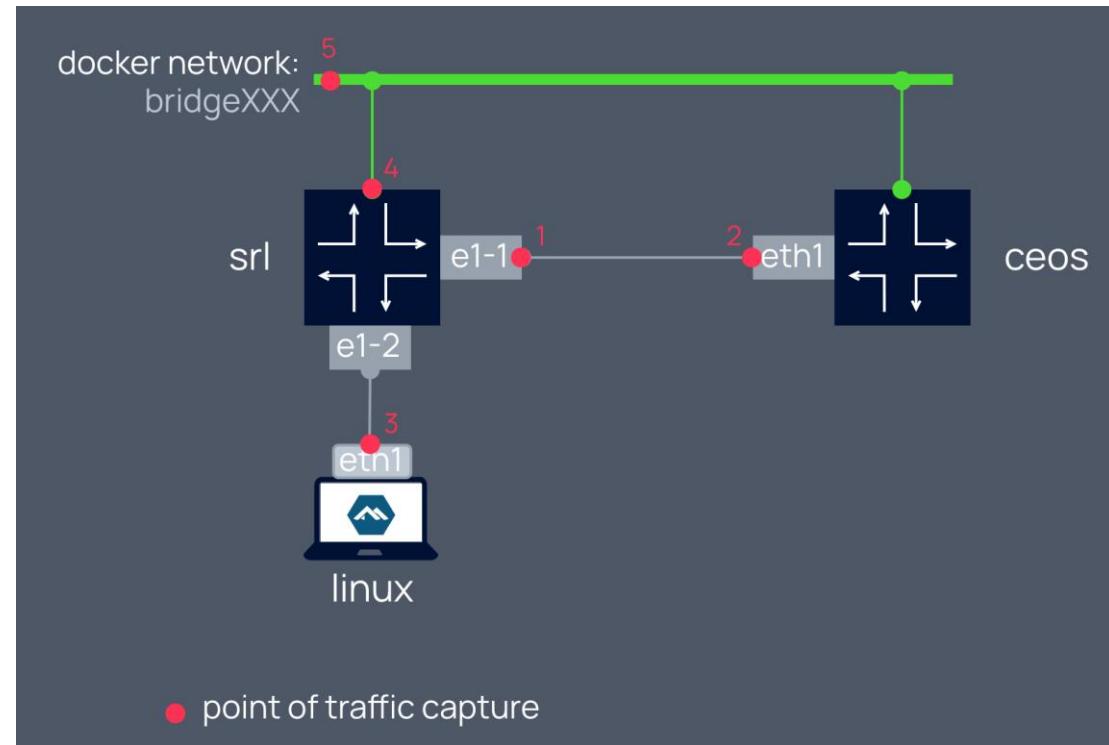
# Packet capture



# Wireshark

Several ways to sniff traffic

- 1 Local capture
  - Via container's shell
- 2 Remote capture
  - Via ssh and pipes
- 3 Web UI capture
  - Via Edgeshark



[containerlab.dev/manual/wireshark](https://containerlab.dev/manual/wireshark)

# Remote capture

A quick way to use Wireshark for packet capture

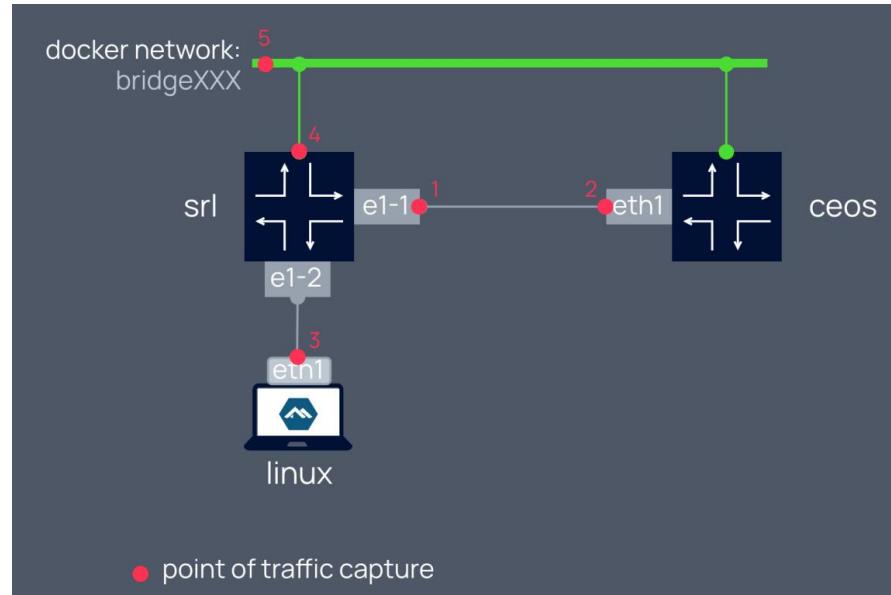
## Command to capture at point #1



```
ssh $clab_host "ip netns exec srl tcpdump  
-U -nni e1-1 -w -" | wireshark -k -i -
```

### Notes:

- Use WSL on windows
- Wrap this long command in a pcap.sh that takes in the note/interface arguments
- No extra packages/modifications required



# Remote capture

A quick way to use Wireshark for packet capture



40-packet-capture

Capturing traffic from **SR OS** port **eth1** (1/1/1) running on **remote** host and displaying in **Wireshark**

```
ssh root@45.76.181.43 \
"sudo ip netns exec clab-vm-sros tcpdump -U -nni eth1 -w -" | \
/mnt/c/Program\ Files/Wireshark/wireshark.exe -k -i -
```



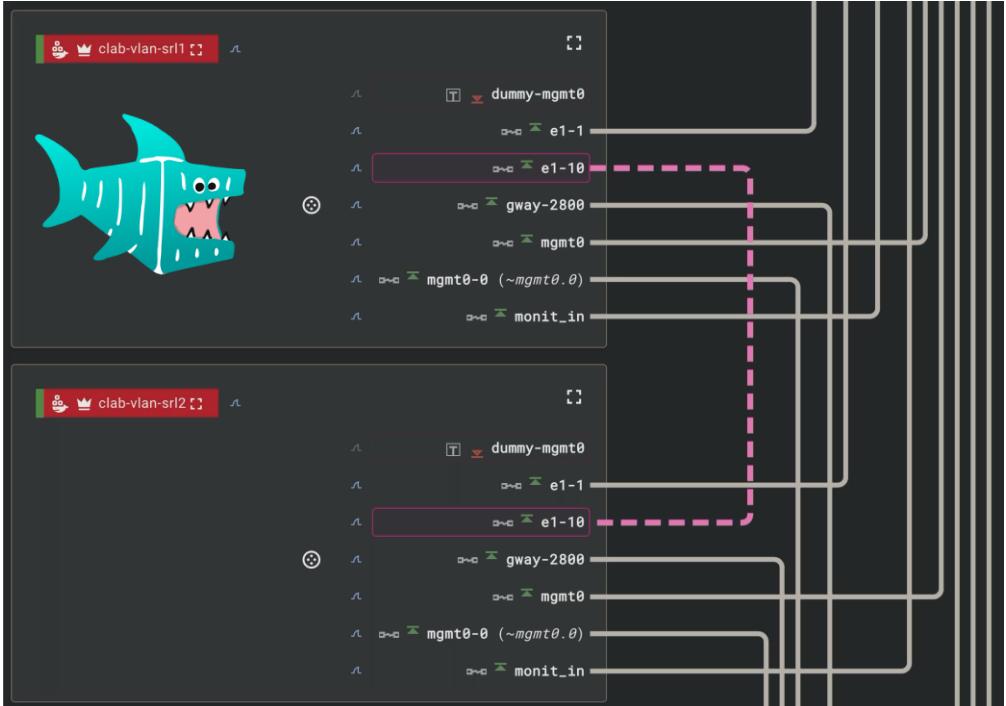
```
ssh root@45.76.181.43 \
"sudo ip netns exec clab-vm-sros tcpdump -U -nni eth1 -w -" | \
/Applications/Wireshark.app/Contents/MacOS/Wireshark -k -i -
```



# Edgeshark

## Web UI for Wireshark

- Way more than this, but that's how we use it
- Web-based, no installation required
- Container-based, no need for dependencies
- Uses native Wireshark
- Open-source
- Free



[containerlab.dev/manual/wireshark/#edgeshark-integration](https://containerlab.dev/manual/wireshark/#edgeshark-integration)

# Deploying Edgeshark



40-packet-capture

## 1 Deploy Edgeshark service

```
● ● ●  
[*]-[<ID>]-[~]  
└─> curl -sL \  
https://github.com/siemens/edgeshark/raw/main/deployments/wget/docker-compose.yaml | docker compose -f - up -d
```

## 2 Install Wireshark capture plugin

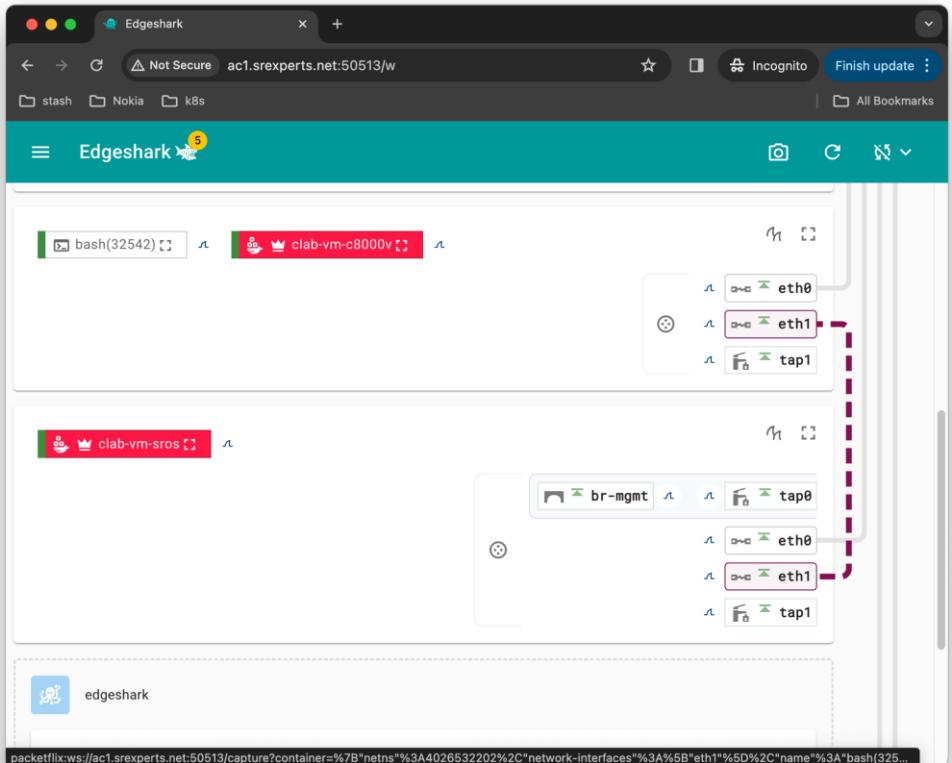
See the workshop's readme!

# Using Edgeshark



<http://45.76.181.43:5001>

When the page opens, hit to refresh the list of containers



# Feature blitz

# Parameter inheritance

Reducing the clutter

```
leaf1:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux  
  
leaf2:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux  
  
leaf3:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux  
  
leaf4:  
  kind: nokia_srlinux  
  image: ghcr.io/nokia/srlinux
```

# Parameter inheritance

## Kind defaults

```
  kinds:  
    nokia_srlinux:  
      image: ghcr.io/nokia/srlinux  
  
    leaf1:  
      → kind: nokia_srlinux  
        image: ghcr.io/nokia/srlinux  
    leaf2:  
      → kind: nokia_srlinux  
        image: ghcr.io/nokia/srlinux  
    leaf3:  
      → kind: nokia_srlinux  
        image: ghcr.io/nokia/srlinux
```

# Parameter inheritance

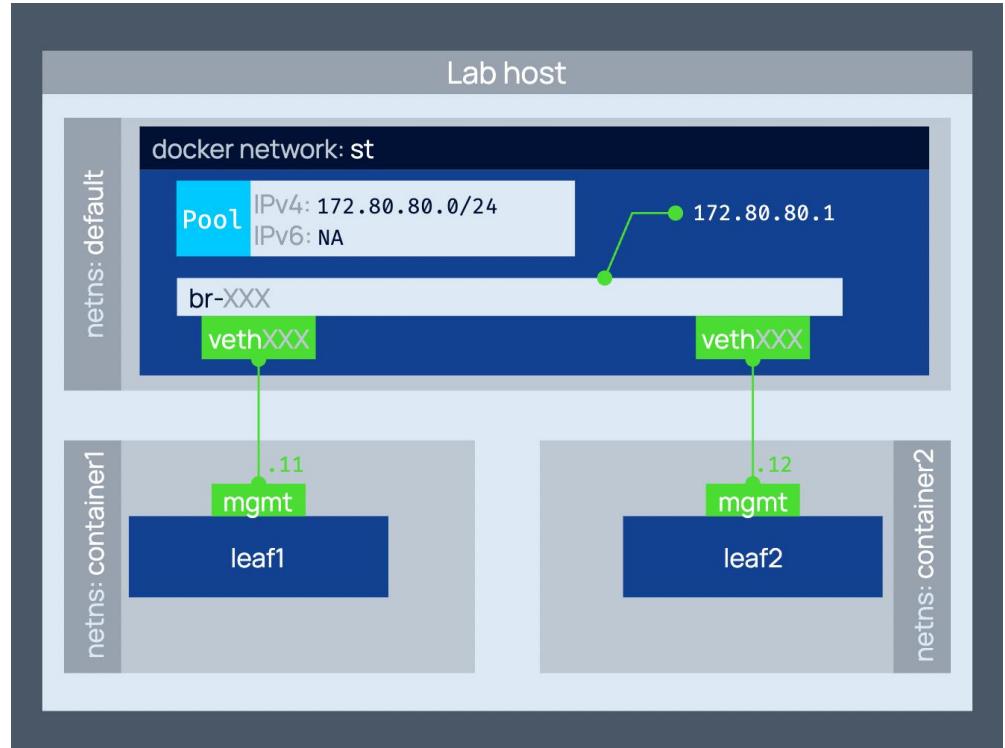
## Global defaults

```
defaults:  
  kind: nokia_srlinux  
  
kinds:  
  nokia_srlinux:  
    image: ghcr.io/nokia/srlinux  
  
leaf1:  
  → kind: nokia_srlinux  
    image: ghcr.io/nokia/srlinux  
  leaf2:  
  → kind: nokia_srlinux  
    image: ghcr.io/nokia/srlinux  
  leaf3:  
  → kind: nokia_srlinux  
    image: ghcr.io/nokia/srlinux
```

# Management network

Statically assigned IP addresses

```
mgmt:  
  network: st  
  ipv4-subnet: 172.80.80.0/24  
  
topology:  
  nodes:  
    leaf1:  
      mgmt-ipv4: 172.80.80.11  
  
    leaf2:  
      mgmt-ipv4: 172.80.80.12
```



# Executing commands

- Exec is a list of commands executed inside the container once it reaches running state
  - configure network params (ip, mac)
  - run the provisioning or workload scripts

```
nodes:  
  server:  
    kind: linux  
    image: alpine:3  
    exec:  
      - ip address add 172.17.0.1/24 dev eth1
```

# Executing commands

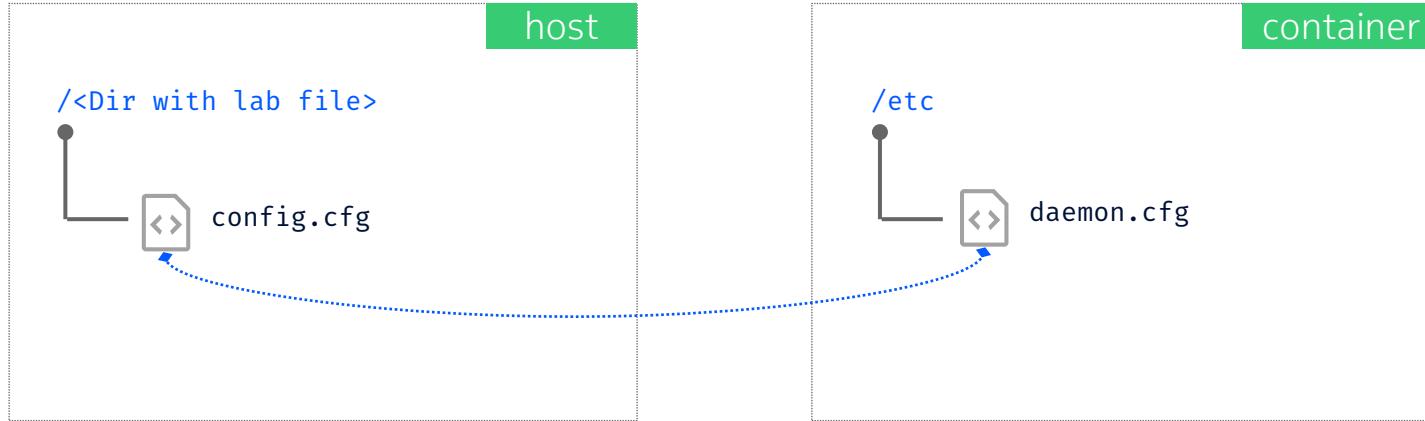
How we used it in the lab?

```
client1:  
  kind: linux  
  exec:  
    - ip address add 172.17.0.1/24 dev eth1  
    - ip -6 address add 2002::172:17:0:1/96 dev eth1  
    - iperf3 -s -p 5201 -D > iperf3_1.log  
    - iperf3 -s -p 5202 -D > iperf3_2.log
```

# File binding

```
server:  
  kind: linux  
  binds:  
    - config.cfg:/etc/daemon.cfg
```

- Bind mount files from host to a container
  - Providing configuration files
  - Providing executable scripts
  - Access to container's files



# File binding

How we used it in the lab?

- Share a config folder with shell scripts from the host to the node
  - Provide iperf.sh script that manages traffic flow

```
client2:  
  kind: linux  
  binds:  
    - configs/client2:/config
```

# Entrypoint and command

- Entrypoint is the “command” that starts in a container
- Command (aka CMD) is a list of arguments passed to the entrypoint

```
server:  
  kind: linux  
  image: alpine:3  
  entrypoint: sleep  
  cmd: "10"
```

# Entrypoint and command

How we used it in the lab?

- Provide configuration options to the processes running in a container

```
gnmic:  
  kind: linux  
  image: ghcr.io/openconfig/gnmic:0.30.0  
  binds:  
    - gnmic-config.yml:/gnmic-config.yml:ro  
  cmd: --config /gnmic-config.yml --log subscribe
```

# Environment variables

- Configure processes via env vars

```
server:  
  kind: linux  
  image: alpine:3  
env:  
  MYVAR:SOMEVALUE
```

# Environment variables

How we used it in the lab?

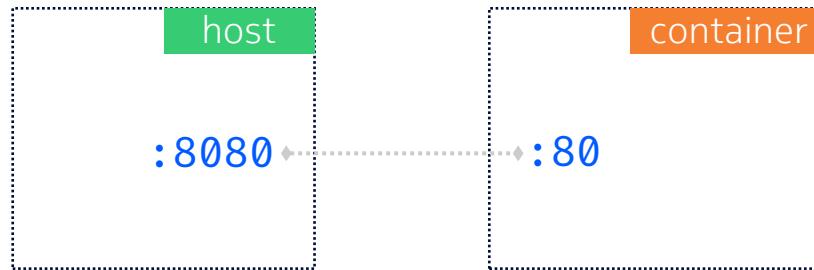
- Provide configuration options for some nodes

```
● ● ●  
grafana:  
  kind: linux  
  image: grafana/grafana:9.5.2  
  env:  
    GF_AUTH_ANONYMOUS_ENABLED: "true"  
    GF_AUTH_ANONYMOUS: "true"
```

# Exposing ports

- Make services inside a container available to containerlab host

```
server:  
  kind: linux  
  image: nginx  
  ports:  
    - 8080:80
```

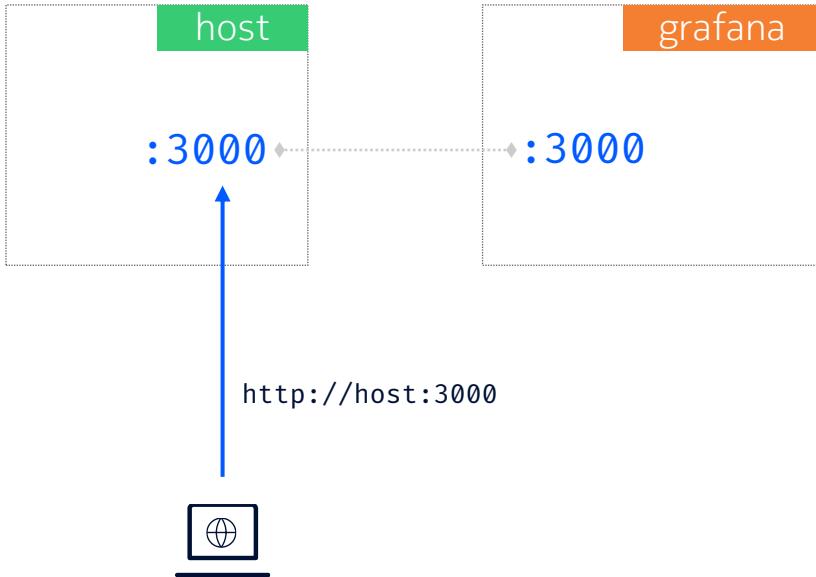


# Exposing ports

How we used it in the lab?

- Expose Grafana Web UI to allow remote access

```
● ● ●  
grafana:  
  kind: linux  
  image: grafana/grafana:9.5.2  
  ports:  
    - 3000:3000
```



# Graphs

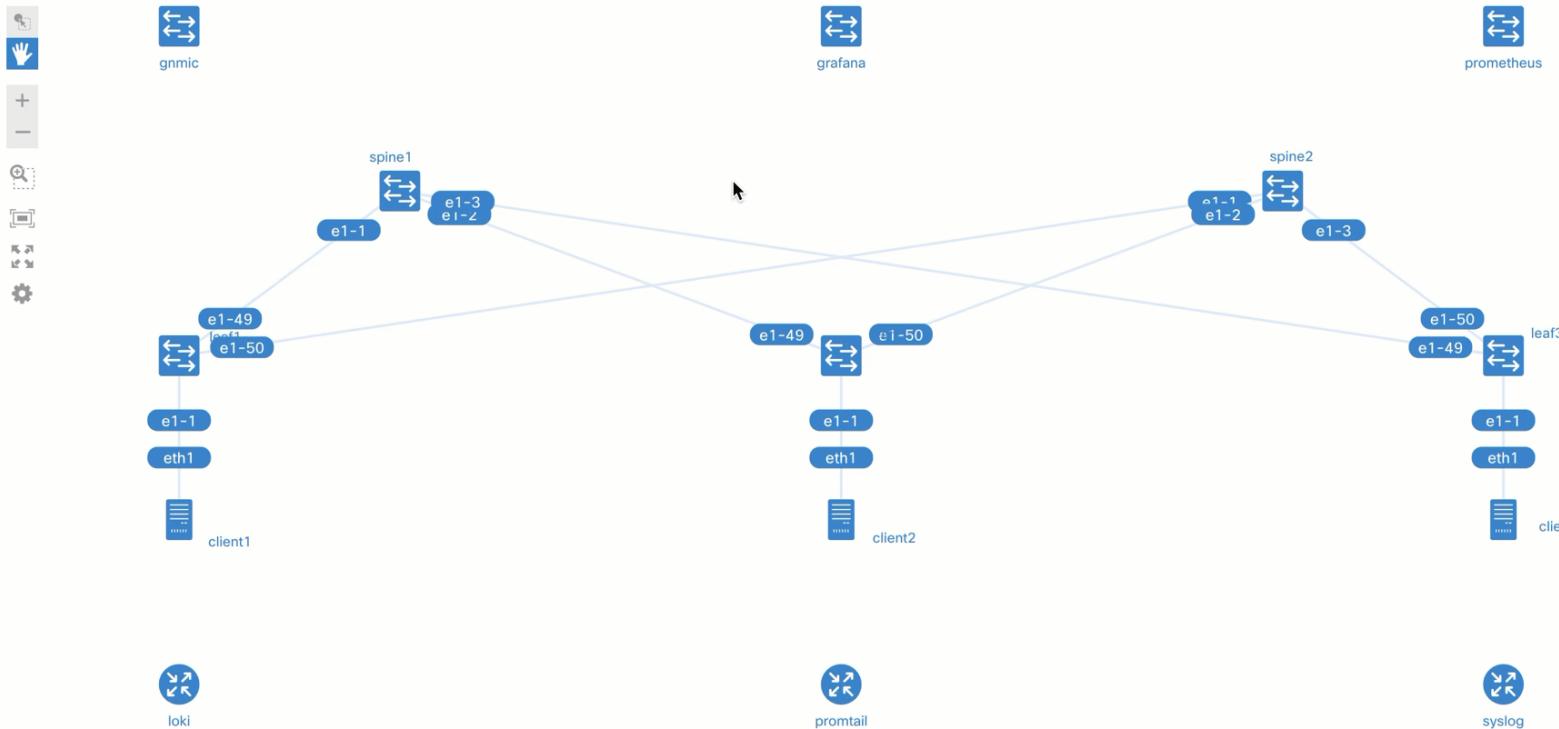
- **graph** command displays the UI with the lab components to visualize the topology
- **group** property may be used to assign nodes a specific role used in sorting the nodes in GUI

```
server:  
  kind: linux  
  image: nginx  
  group: server
```

# ContainerLab Topology ST

Horizontal Layout

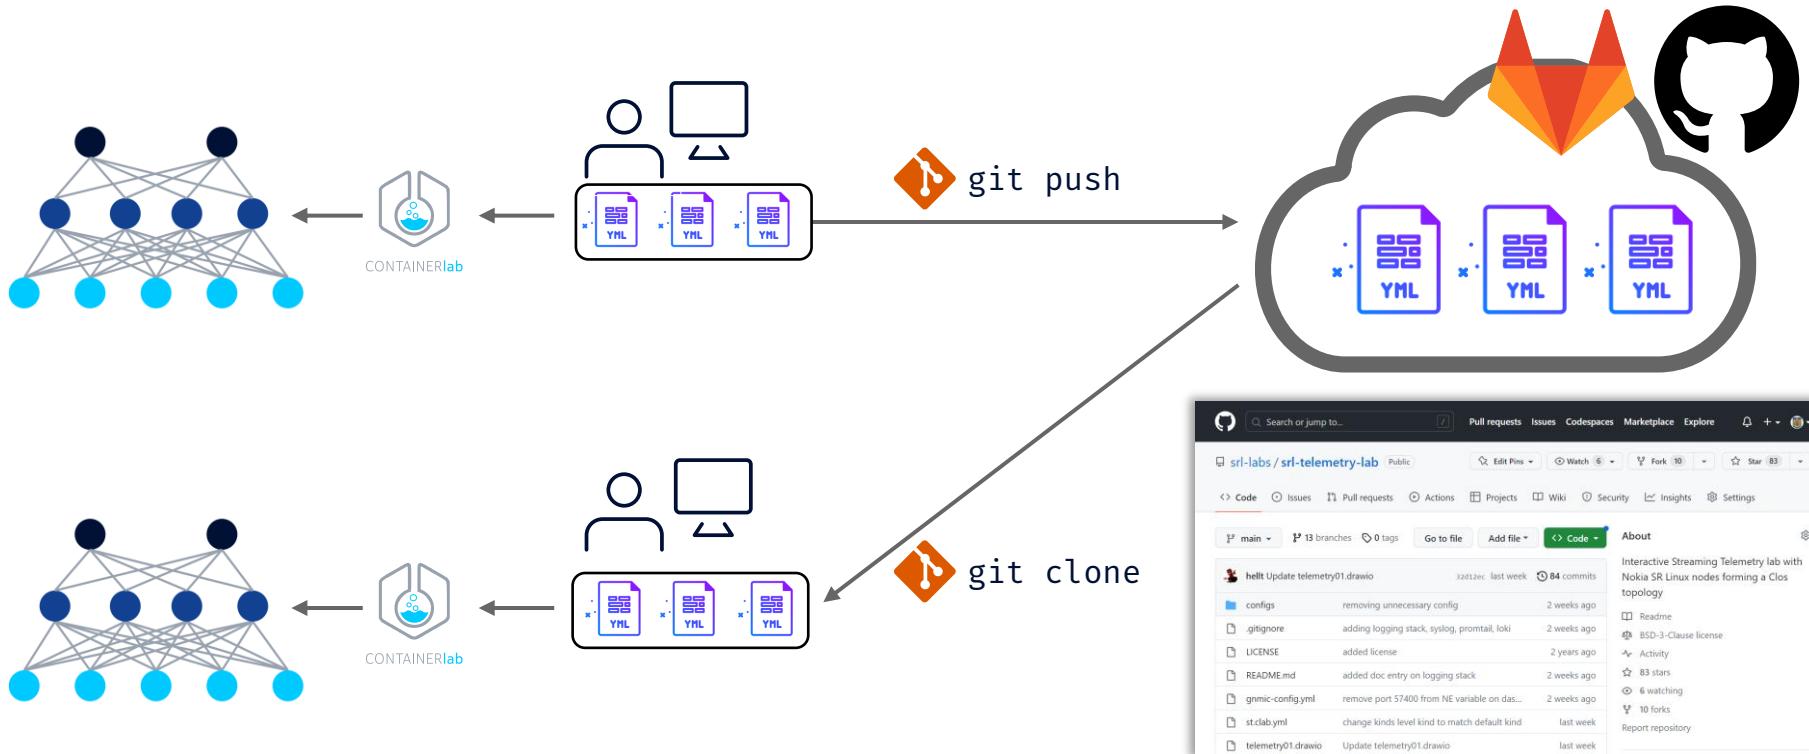
Vertical Layout



# Sharing and finding containerlabs

# Share your labs

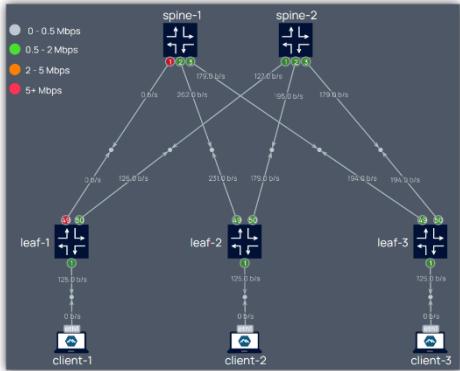
## Lab As Code



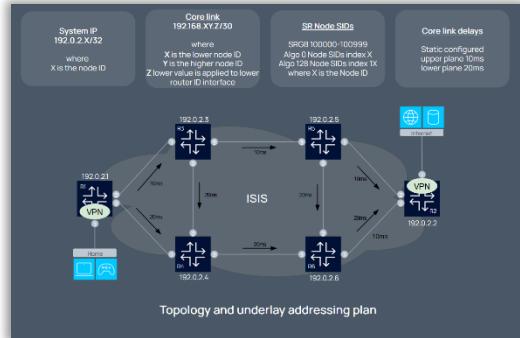
# Lab examples



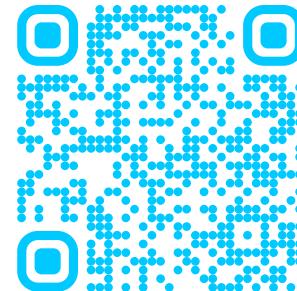
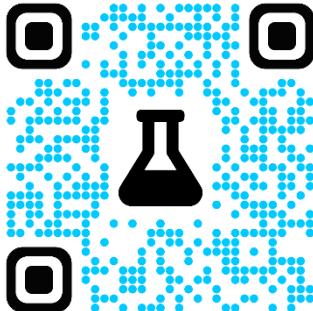
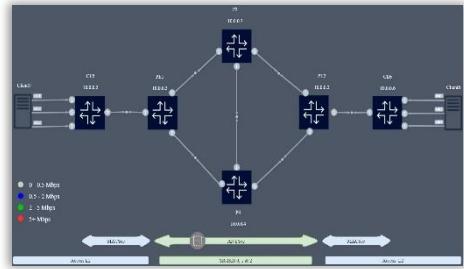
## Streaming Telemetry



## Segment Routing



## AnySec/MACSec

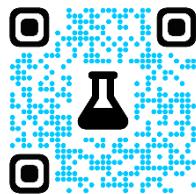
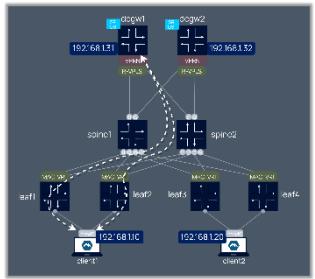


NOKIA

# More labs!



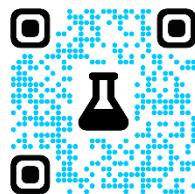
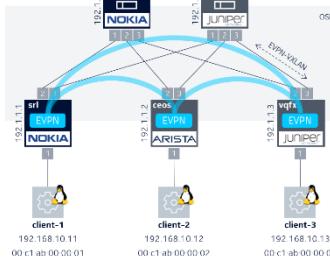
## Nokia EVPN Interop



[srl-labs/nokia-evpn-lab](https://github.com/srl-labs/nokia-evpn-lab)



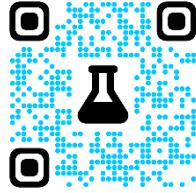
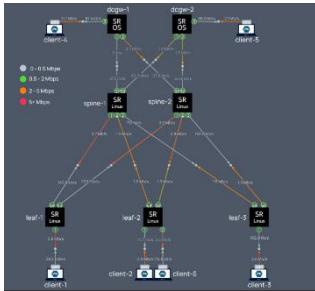
## Multivendor EVPN



[srl-labs/multivendor-evpn-lab](https://github.com/srl-labs/multivendor-evpn-lab)



## SR Linux & SROS Telemetry



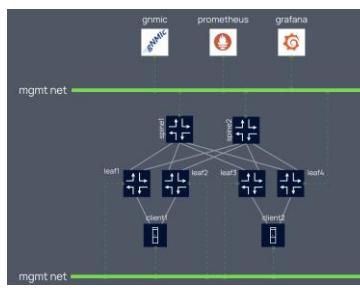
[srl-labs/srl-sros-telemetry-lab](https://github.com/srl-labs/srl-sros-telemetry-lab)



CONTAINERlab



## SR Linux Oper-Group



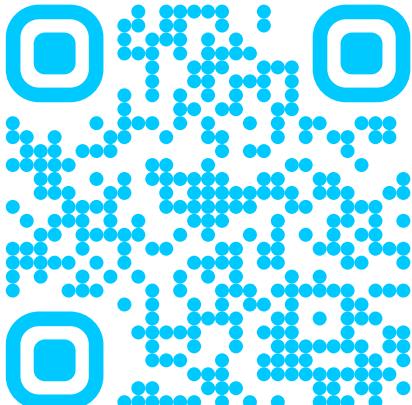
[srl-labs/opergroup-lab](https://github.com/srl-labs/opergroup-lab)

NOKIA

# Distributed collection of containerlabs

Add [clab-topo](#) topic to your repo

- Make your labs easily discoverable in a distributed way
- Each repo represents a runnable topology
- Publish your lab and others will see it!



The screenshot shows the GitHub Topics page for the 'clab-topo' topic. The top navigation bar includes 'Explore', 'Topics' (which is underlined), 'Trending', 'Collections', 'Events', and 'GitHub Sponsors'. A search bar and various icons are also present. Below the header, the '# clab-topo' section is shown with a 'Star' button. A message states 'Here are 20 public repositories matching this topic...'. Two repository cards are visible: 'srl-labs / srl-telemetry-lab' and 'srl-labs / nokia-segment-routing-lab'. Both cards include details like star count, issues, pull requests, and a brief description. The bottom right corner features a dark overlay with the word 'About' and a description of the lab.

Explore Topics Trending Collections Events GitHub Sponsors

# clab-topo Star

Here are 20 public repositories matching this topic...

Language: All Sort: Most stars

srl-labs / srl-telemetry-lab Starred 110

Code Issues Pull requests

Interactive Streaming Telemetry lab with Nokia SR Linux nodes forming a Clos topology  
clab-topo

Updated on Nov 19, 2023 Shell

srl-labs / nokia-segment-routing-lab

Code Issues Pull requests

A lab to demonstrate Flex-Algo use case  
clab-topo

Updated on May 24, 2023 JavaScript

About

Interactive Streaming Telemetry lab with Nokia SR Linux nodes forming a Clos topology

gnmi streaming-telemetry clab-topo

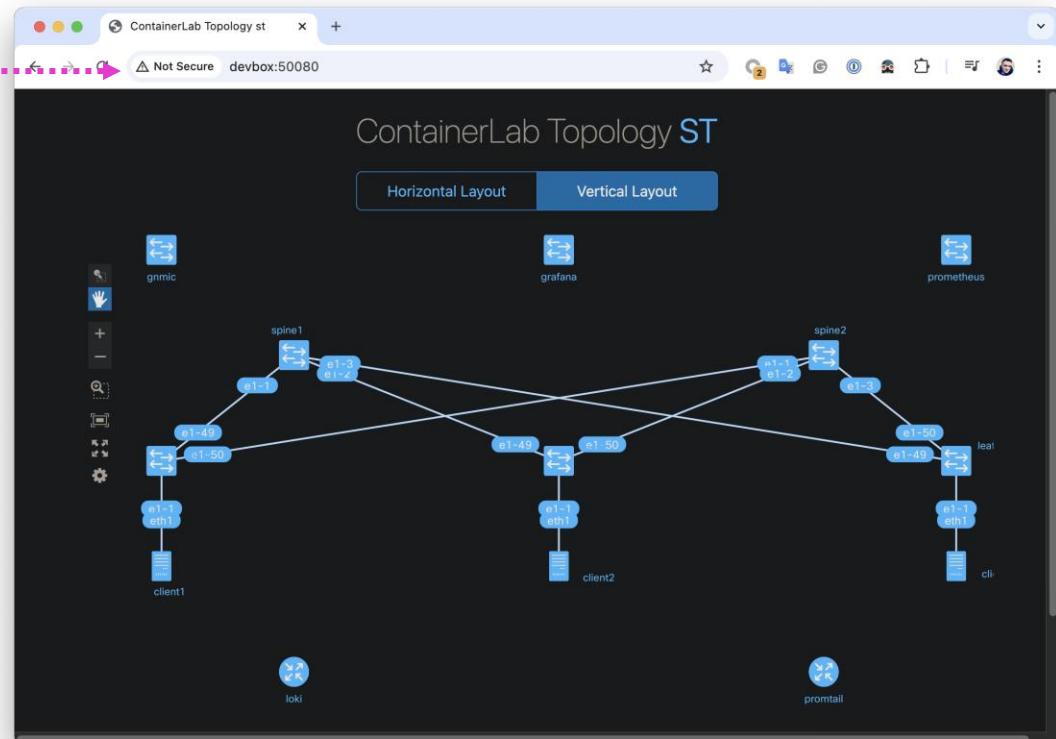
# What exactly we just deployed?



```
sudo clab graph
```

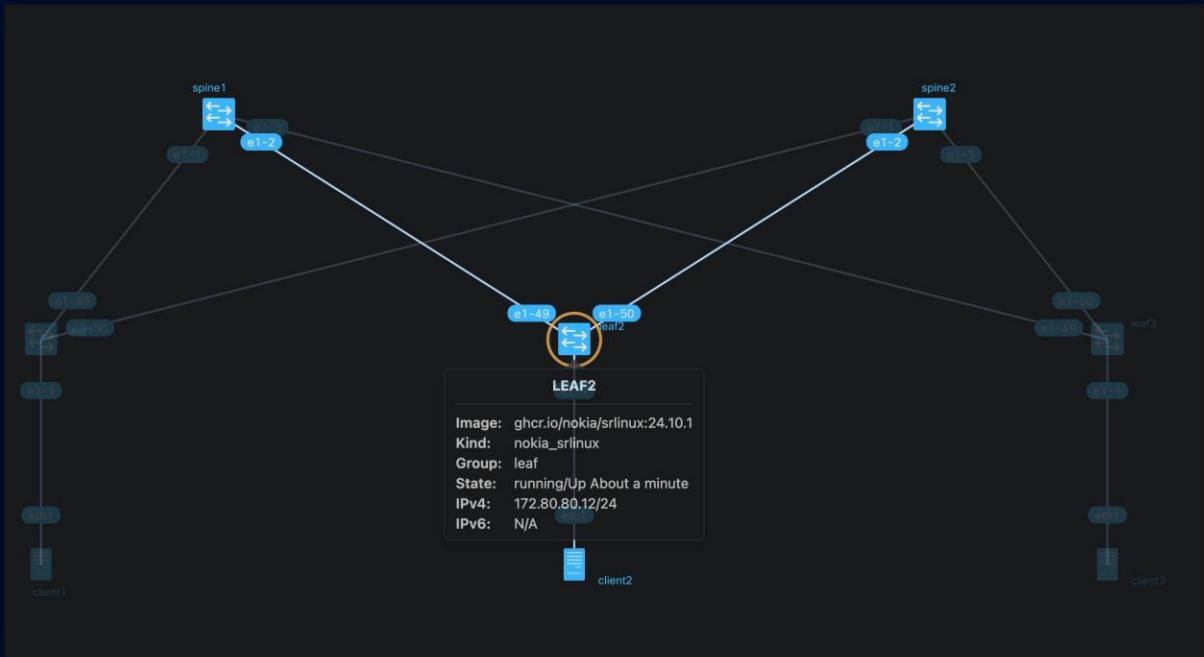
```
INFO[0000] Parsing & checking topology file:  
st.clab.yml  
INFO[0000] Serving topology graph on  
http://0.0.0.0:50080
```

<http://iID.innog8.net:50080>



# Interactive graph

- Display node and link info
  - Highlight the connected paths
  - Show mgmt IP info
  - Works with deployed and not deployed topologies
- 
- Not editable
  - Opinionated
  - Only one "pane" - physical



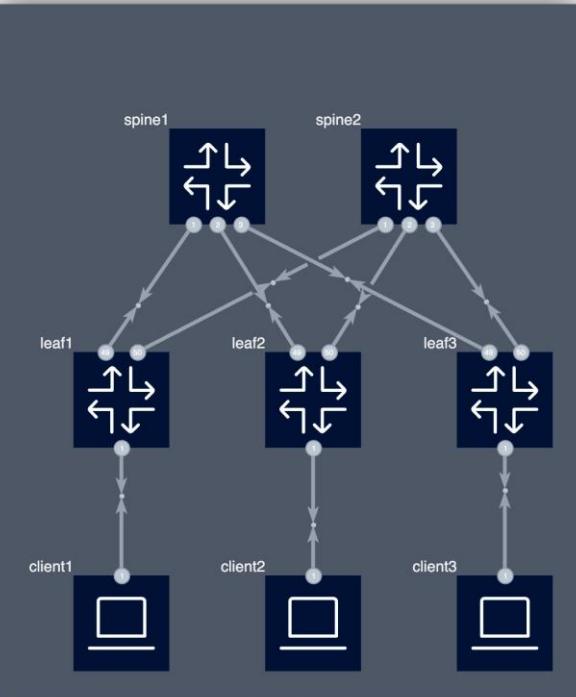
# clab-io-draw

More diagrams to the god of diagrams



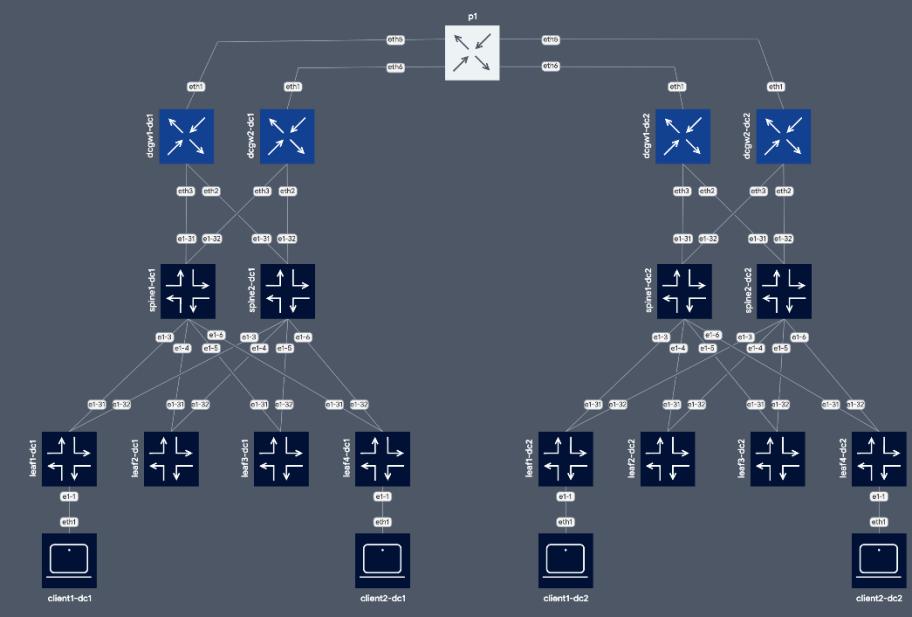
```
sudo clab graph --drawio  
--drawio-args "--theme grafana_dark"
```

- Generates the real “drawio” file
- Editable, shareable, customizable
- Supports themes
- Generates flow panel configuration
- Interactive dialog to customize icons positioning, etc



# clab-io-draw

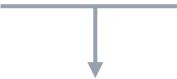
- A perfect start for a beautiful documentation diagram
- Explore projects' homepage to check the available options
- Create your own style and contribute to the collection of the color themes



NOKIA

# Graph options

GRAPH



Web graph



Drawio



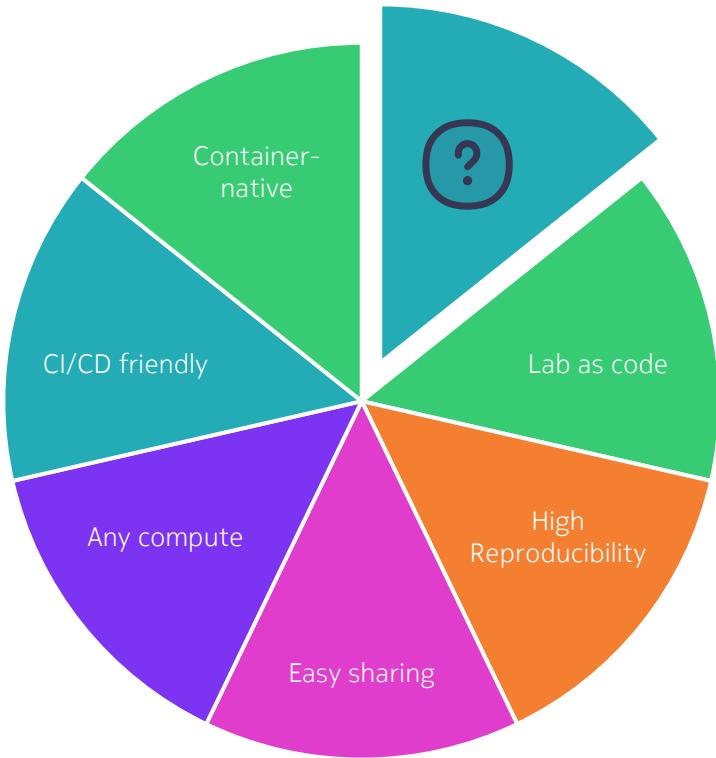
Mermaid



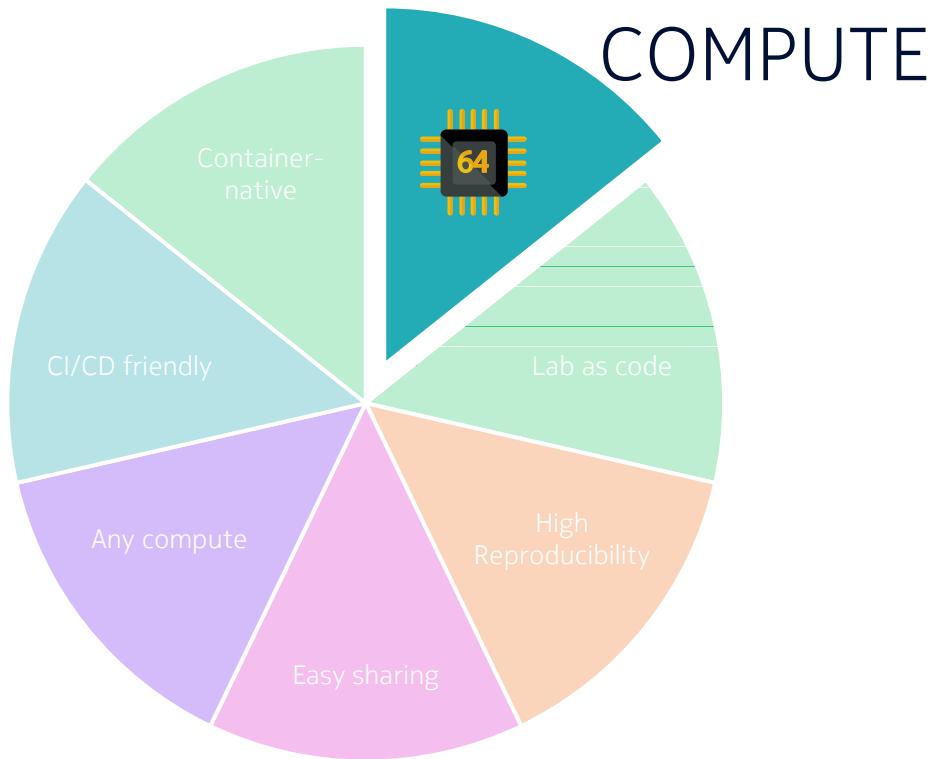
[containerlab.dev/cmd/graph/](https://containerlab.dev/cmd/graph/)

# One more thing...

# A missing piece



# A missing piece



# Codespaces

Free compute\*

- Microsoft-backed compute cloud
- Generous free tier with a monthly quota reset
- 120cpu/hours a month **FREE**
- Containerlab integrated with Codespaces
- No \$\$\$ until explicitly committed

## Nokia SR Linux Streaming Telemetry Lab

SR Linux has first-class Streaming Telemetry support thanks to [100% YANG coverage](#) of state and config data. The holistic coverage enables SR Linux users to stream **any** data off of the NOS with on-change, sample, or target-defined support. A discrepancy in visibility across APIs is not about SR Linux.

This lab represents a small Clos fabric with [Nokia SR Linux](#) switches running as containers. The lab topology consists of a Clos topology, plus a Streaming Telemetry stack comprised of [gnmic](#), [prometheus](#) and [grafana](#) applications.



### Run In Codespaces

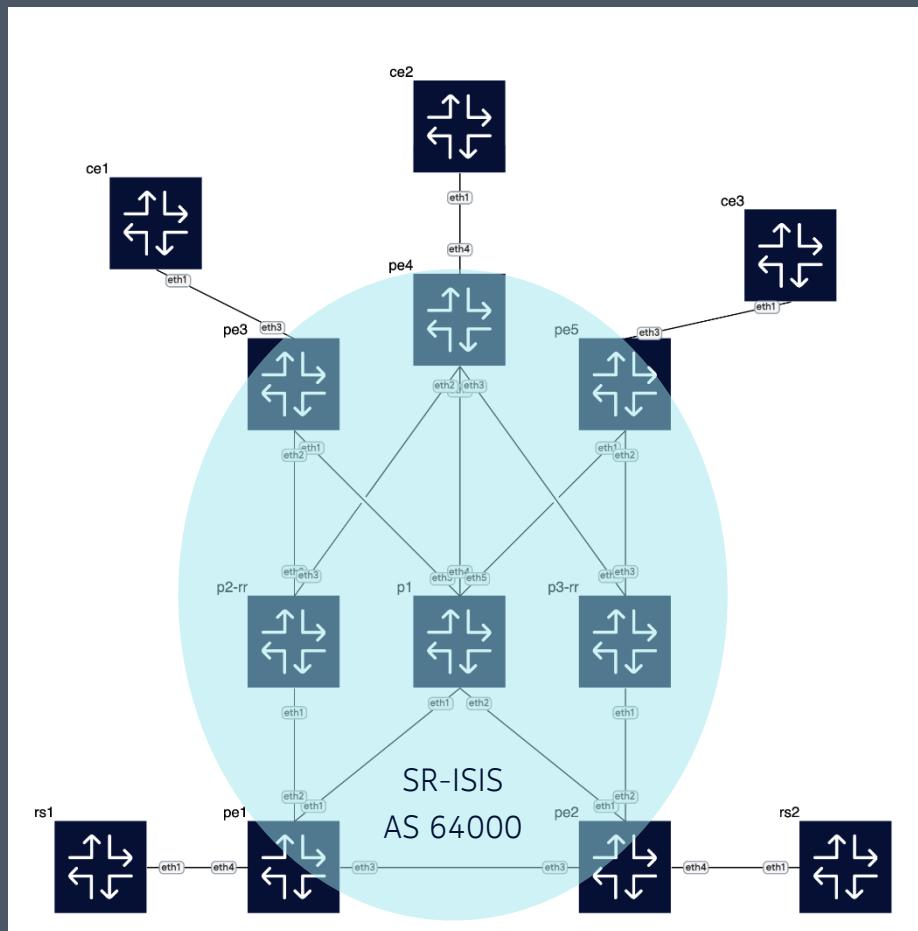
[Run this lab in GitHub Codespaces for free.](#)

[Learn more](#) about Containerlab for Codespaces.

# IXP Peering Lab

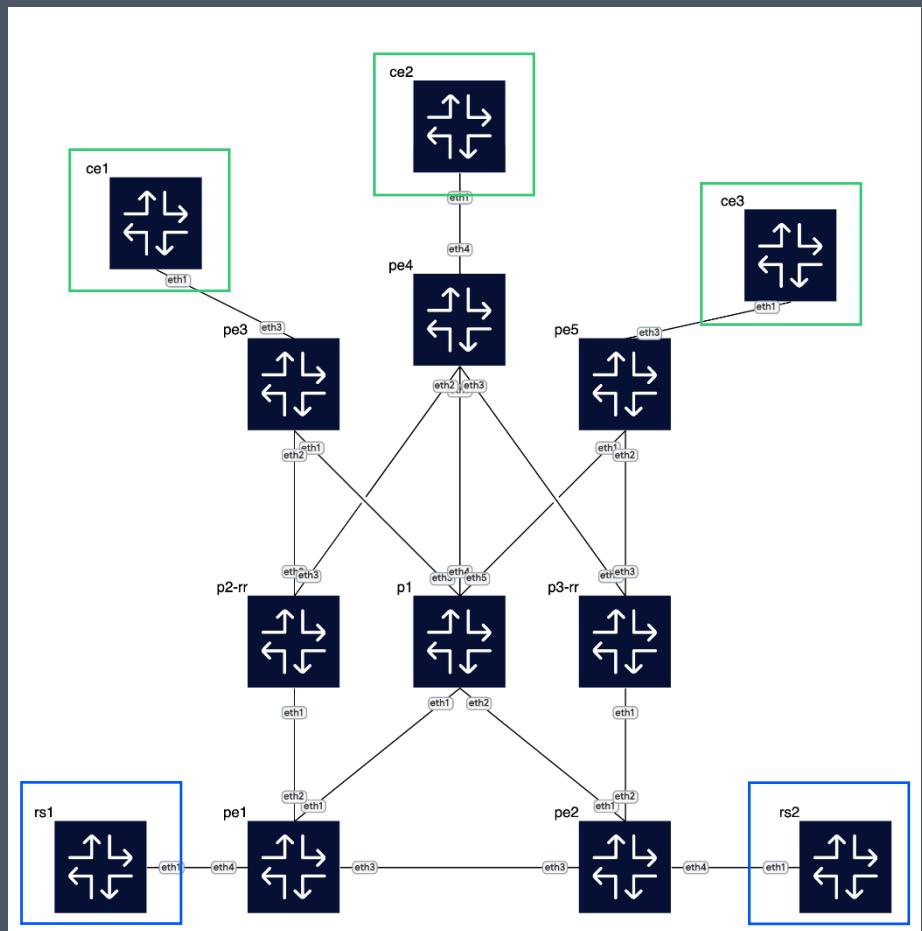
# IXP Peering lab (1/2)

- This lab is a representation of a modern IXP peering network using an EVPN service architecture
- The current architecture also focuses on SR-MPLS (SR-ISIS) for the underlay, with BGP providing EVPN service route exchanges over this underlay.



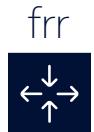
## IXP Peering lab (2/2)

- BIRD routing daemon as a Route Server is used (rs1), as well as OpenBGPd (rs2) for diversity and sustainability reasons.
- The CEs are IXP clients based on FRR, which have peering towards the Route Servers and learn other CE prefixes through them



# FRR node as CE role

## Purpose



- FRR configuration is split between the two files:
  - frr.conf - contains the basic FRR configuration, which includes the most simple BGP configuration to enable peering.
  - daemons.cfg - contains the list of FRR daemons to be started

# FRR node

## Configuration

frr



- FRR configuration is split between the two files:
  - frr.conf - contains the basic FRR configuration, which includes the most simple BGP configuration to enable peering.
  - daemons.cfg - contains the list of FRR daemons to be started

# OpenBGPD node as Route Server role

## Purpose

- Almost every Internet eXchange Point (IXP) leverages a Router Server (RS) to simplify peering between members of the exchange who exercise an open policy peering.
- A Route Server is a software component connected to the IXP network which acts as a BGP speaker with whom members peer to receive BGP updates from each other.
- Lots of IXPs are interested in introducing OpenBGPD as a second Route Server in their networks and this lab opens the doors to explore "OpenBGPD as a Route Server" use case.

OpenBGPD



# OpenBGPd node

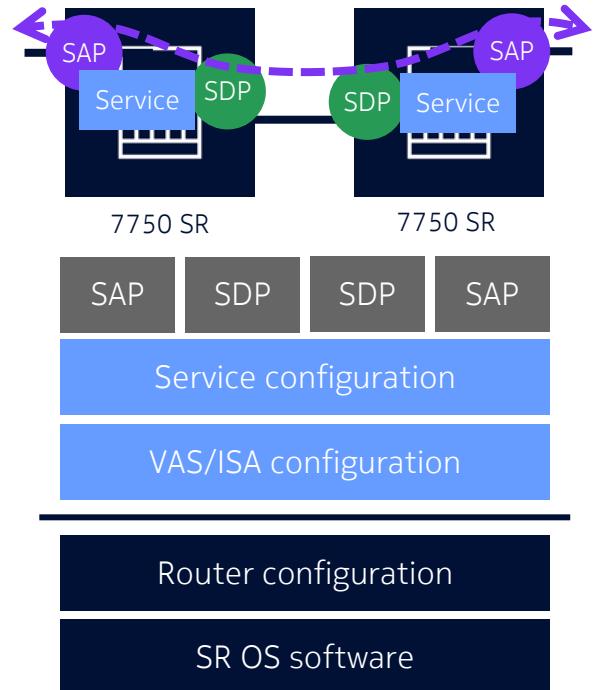
## Configuration

- To provide OpenBGPd with a configuration file we leverage the binds property that works exactly like bind mount in Docker.
- You specify a source file path and the corresponding path inside the container process. For openbgpd node we take the openbgpd.conf file and mount it inside the container by the /etc/bgpd/bgpd.conf path. This will make OpenBGPd to read this config when the process starts.
- One last step left for the openbgpd node: to configure the eth1 interface that connects the Route Server to the IXP LAN. Links setup is covered in details in the next chapter, but, for now, just keep in mind that we can configure interfaces of a container using the exec option and ip utility..

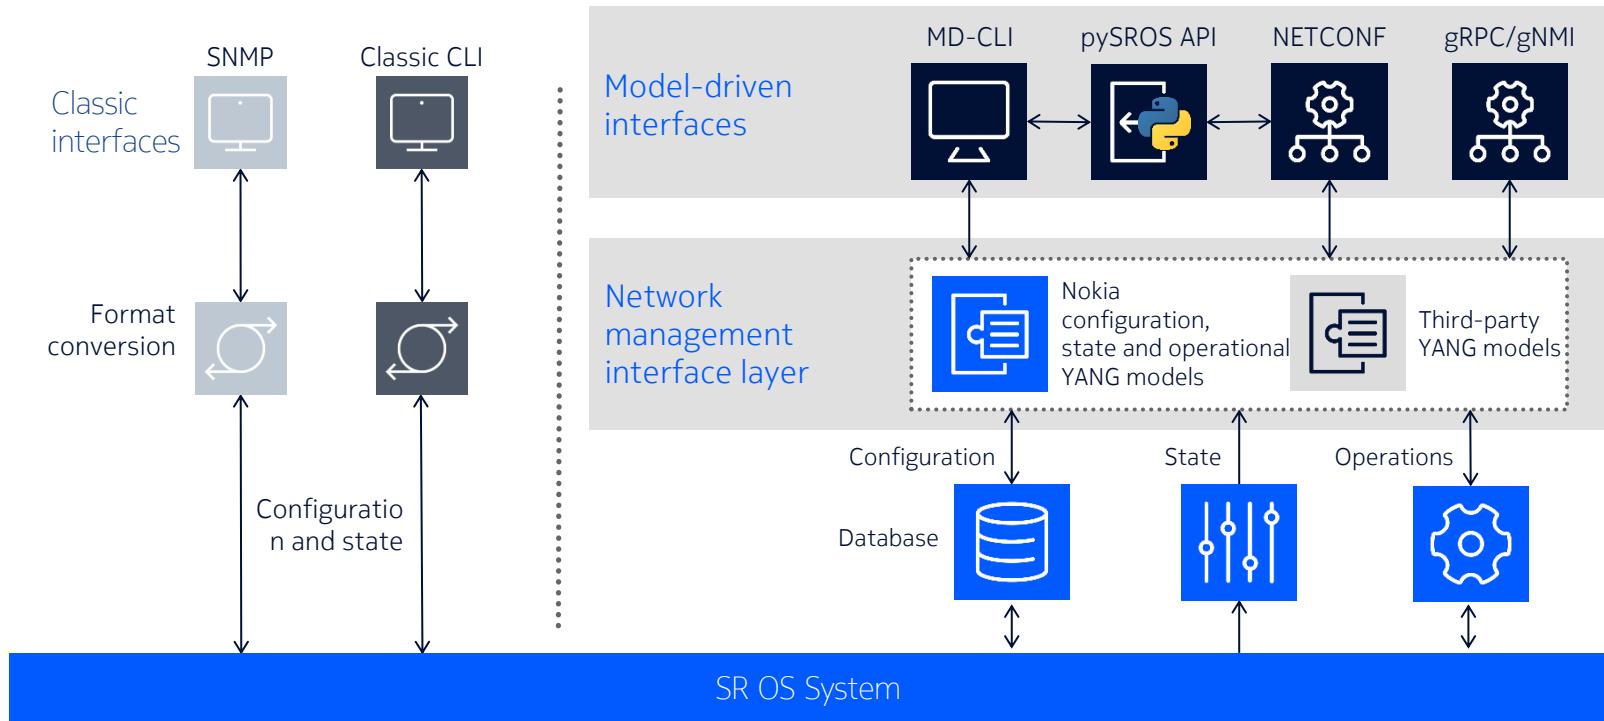


# SR OS services architecture

- SR OS, the “Service Router Operating System”, was designed to support services from the very beginning
- Service model uses logical service entities to construct a service
  - Provides a uniform, service-centric configuration, management, and billing model
  - Service access points (SAPs) abstract the service interface from the physical interface and type
  - Service distribution points (SDPs) abstract the network path into a logical tunnel that interconnect services
- Separate service configuration and router configuration enables different access levels to service configuration
- Strict security controls protect the router configuration and allow granular customizable user access to CLI commands
- QoS policies, filter policies, and accounting policies are applied to each service which simplifies provisioning

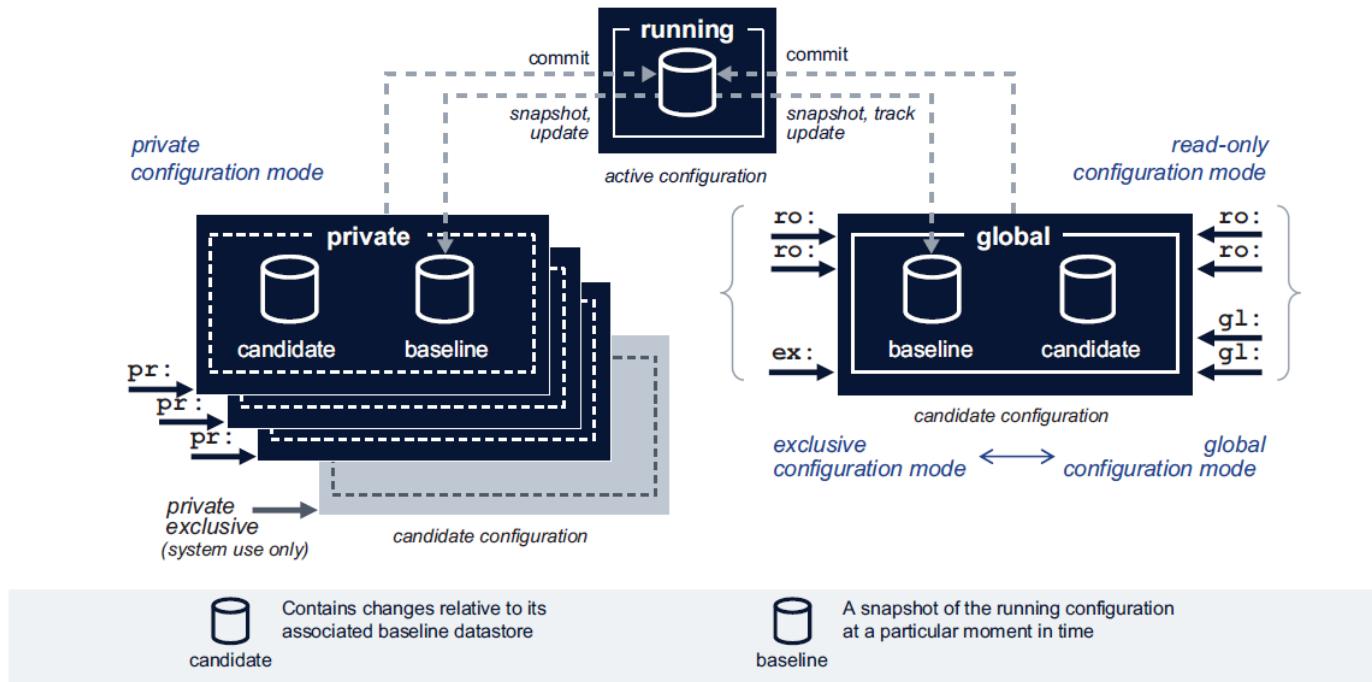


# SR OS management infrastructure architecture



# Model Driven CLI – Transactional Configuration

Accessing different datastores through configuration modes



sw0637

# Transactional configuration management commands

- **info**: shows configuration of current datastore
- **commit [confirmed]**: commit changes to the running datastore
- **compare**: compare datastore changes
- **delete**: remove explicit configuration and sets to default
- **discard**: discard changes in the candidate datastore
- **load**: load a file to the candidate datastore
- **rollback**: rollback to a previous configuration in the candidate
- **save**: save the candidate datastore to a file
- **validate**: validate changes made to the candidate datastore

# MD-CLI Prompt

- The MD-CLI has a two-line prompt
- First line shows dynamic state and current context information
  - Full context display with all variables expanded, no abbreviations and intelligent wrapping
  - Context can be copied and pasted
  - Context always shows exactly where the operator was, useful for questions and support cases
  - Dynamic state shows uncommitted changes indicator (\*), and the configuration mode (gl: global, ex: exclusive, ro: read-only, pr: private)
- Second line shows static user and system information in a format that stays the same
  - Easier to screen scrape for automation
- Leading newline is always printed before a new prompt, but can be configured not to

# MD-CLI Prompt example

Configuration mode (gl: global, ex: exclusive, ro: read-only, pr: private)

Active CPM slot

Uncommitted changes indicator (\*)

Present working context

```
(ex) [configure]
A:admin@cses-V93# router

(ex) [configure router "Base"]
A:admin@cses-V93# interface foo

(ex) [configure router "Base" interface "foo"]
A:admin@cses-V93# ipv6 address 2406:c800:a1ca:7e1:7950:7750:7210:7705

*(ex) [configure router "Base" interface "foo" ipv6 address 2406:c800:a1ca:
7e1:7950:7750:7210:7705]
A:admin@cses-V93# █
```

# MD-CLI – Cheat sheet

- “info” show non-default configuration in current location, not in operational datastore
- “info detail” show default non-default configuration in current location
- “admin show configuration” show all running configuration
- “tree (flat)” show configuration tree from current location
- “back” go exactly 1 level back in configuration tree
- “exit” exit current section in configuration tree
- “exit all” go to root of configuration tree
- “... | match <string>” filter output on occurrence of <string> - multiple options possible
- “admin save” save running configuration to location as configured in BOF file
- “admin save <L>” save running configuration to local or remote location <L>
- “admin reboot” hard reboot of router
- “admin support-mode” Enable the shell and kernel commands

# Chassis Slots and Cards

## Port Types

Router ports must be configured as either access, hybrid, or network. The default is network.

- **Access ports** — Configured for customer facing traffic on which services are configured. If a Service Access Port (SAP) is to be configured on the port or channel, it must be configured as an access port or channel. When a port is configured for access mode, the appropriate encapsulation type must be configured to distinguish the services on the port or channel. Once a port has been configured for access mode, one or more services can be configured on the port or channel depending on the encapsulation value.
- **Network ports** — Configured for network-facing traffic. These ports participate in the service provider transport or infrastructure network. Dot1q is supported on network ports.
- **Hybrid ports** — Configured for access and network-facing traffic. While the default mode of an Ethernet port remains network, the mode of a port cannot be changed between the access, network, and hybrid values unless the port is shut down and the configured SAPs or interfaces are deleted. Hybrid ports allow a single port to operate in both access and network modes.

# Chassis Slots and Cards

## Configuring Connectors and Connector Ports

Some assemblies have support for QSFP28 or QSFP-DD transceiver modules. These modules have different variants, some of which provide multiple physical ports out of a single module (breakout modules). There is a QSFP28 breakout module that supports ten physical 10 Gb Ethernet ports. On assemblies that support these breakout variants, the front panel cages are modeled as connectors rather than as direct ports. The connector must be configured for the type of breakout module that is to be inserted and then the appropriate ports are created and can be configured.

The connector reference is in the format *slot/mda/connector* (for example, 1/1/c3) and the ports owned by the connector use the format *slot/mda/connector/port*. For example, in a 7750 SR-1 with the 6-port QSFP28 mda-e-xp installed in the first MDA slot, initially there are no ports available, only six connectors:

```
show mda
```

```
=====
```

```
MDA Summary
```

```
=====
```

Slot	Mda	Provisioned Type	Admin	
Operational		Equipped Type (if different)	State	State
1	1	s36-100gb-qfp28:he2400g+	up	up

```
show port
```

```
=====
```

```
Ports on Slot 1
```

```
=====
```

```
Port Admin Link Port Cfg Oper LAG/ Port Port Port C/QS/S/XFP/  
Id State State MTU MTU Bndl Mode Encp Type MDIMDX
```

```
=====
```

```
1/1/c1 Up No Down - unkn unkn conn
```

```
1/1/c2 Up No Down - unkn unkn conn
```

```
1/1/c3 Up No Down - unkn unkn conn
```

```
1/1/c4 Up No Down - unkn unkn conn
```

```
1/1/c5 Up No Down - unkn unkn conn
```

```
1/1/c6 Up No Down - unkn unkn conn
```

# Chassis Slots and Cards

## Configuring Connectors and Connector Ports

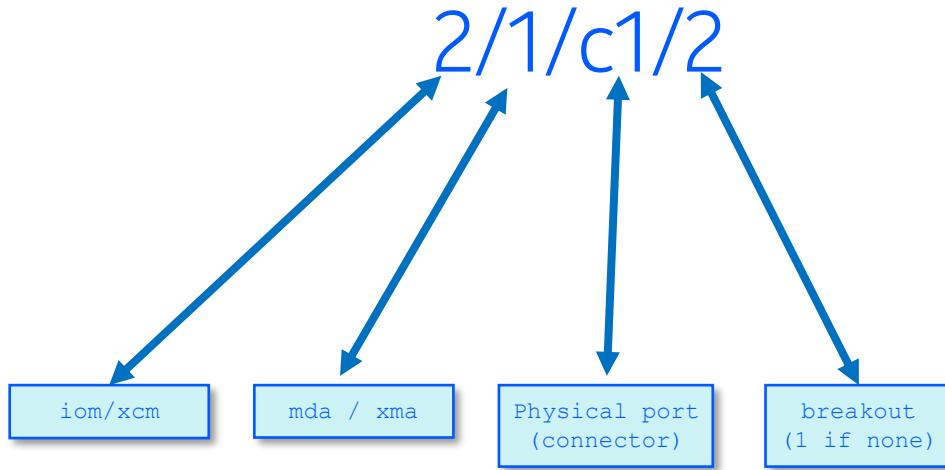
After configuring a module with four 10 Gb breakout ports in connector position 1/1/c1 and a module with one 100 Gb breakout port in connector position 1/1/c2, the physical ports are created

```
configure {
    port 1/1/c1 {
        admin-state enable
        connector {
            breakout c4-10g
        }
    }
    port 1/1/c2 {
        admin-state enable
        connector {
            breakout c1-100g
        }
    }
}
```

```
show port
=====
Ports on Slot 1
=====
Port Admin Link Port Cfg Oper LAG/ Port Port Port C/QS/S/XFP/
Id State State MTU MTU Bndl Mode Encp Type MDIMDX
-----
1/1/c1 Up No Link Up - unkn unkn conn 40GBASE-SR4
1/1/c1/1 Down No Down 9212 9212 - netw null xgige
1/1/c1/2 Down No Down 9212 9212 - netw null xgige
1/1/c1/3 Down No Down 9212 9212 - netw null xgige
1/1/c1/4 Down No Down 9212 9212 - netw null xgige
1/1/c2 Down No Link Up - unkn unkn conn 100GBASE-LR4
1/1/c2/1 Down No Down 1578 1578 - netw null cgige
1/1/c3 Up No Down - unkn unkn conn
1/1/c4 Up No Down - unkn unkn conn
1/1/c5 Up No Down - unkn unkn conn
1/1/c6 Up No Down - unkn unkn conn
```

# Chassis Slots and Cards

## Port labelling



# Interfaces overview

- "router interface" in SROS is a logical interface that can be either loopback or network interface
- "router interface" called "system" exists by default and is used as source for all routing and mgmt
- Additional loopback interfaces can be configured
- "router interface" is identified by a name (32 characters)
- Description can be added to router interface (160 characters)
- Minimum configuration for a network interface:
  - IPv4 address
  - Port or LAG + optionally VLAN
- Interfaces can be applied to protocols (e.g. OSPF, LDP,..)

```
configure {
    router "Base" {
        interface "system" {
            ipv4 {
                primary {
                    address 1.1.1.3
                    prefix-length 32
                }
            }
        }
    }
}
```

```
configure {
    router "Base" {
        interface "int-sr-b-lag2" {
            admin-state enable
            port lag-2
            ipv4 {
                primary {
                    address 192.168.2.2
                    prefix-length 30
                }
            }
        }
    }
}
```

# Interfaces hierarchy

```
port 1/1/c1/1 {
    admin-state enable
    ethernet {
        mode network
        encap-type null
        mtu 9212
    }
}
```

```
router "Base" {
    interface "int-sr-b-lag2" {
        admin-state enable
        port 1/1/c1/1
        ipv4 {
            primary {
                address 192.168.2.2
                prefix-length 30
            }
        }
    }
}
```



# Segment Routing MPLS

- Configuration steps for ISIS example:
  1. Create a label range for sr-labels from dynamic label range
    - Configured range is the Segment Routing Global Block (SRGB)
  2. Advertise to neighbor that router is SR-capable & Define flooding scope (area) for ISIS
  3. Node SID label is configured from SRGB
  4. Define Prefix SID range
    - “global” – SID defined from SRGB
  5. Enable Segment Routing

```
(g1) [/configure router "Base"]
A:admin@R1# info
-- SNIP --
mpls-labels {
    sr-labels start 400000 end 401000
}
...
(g1) [/configure router "Base" isis 0]
A:admin@R1# info
-- SNIP --
    traffic-engineering true
    advertise-router-capability area
    interface "system" {
        interface-type point-to-point
        ipv4-node-sid {
            index 101
        }
    }
    segment-routing {
        admin-state enable
        prefix-sid-range {
            global
        }
    }
}
```

1  
2  
3  
4  
5

# Routing protocols - BGP configuration (GRT)

```
router "Base" {
    autonomous-system 64400
    bgp {
        rapid-withdrawal true
        peer-ip-tracking true
        group "IXP-iBGP" {
            type internal
            family {
                evpn true
            }
            import {
                policy ["from-as-64400"]
            }
            export {
                policy ["export-prefixes"]
            }
        }
        neighbor "10.10.10.2" {
            group "IXP-iBGP"
        }
        neighbor "10.10.10.3" {
            group "IXP-iBGP"
        }
        neighbor "10.10.10.4" {
            group "IXP-iBGP"
        }
        neighbor "10.10.10.5" {
            group "IXP-iBGP"
        }
    }
}
```

[ipv4] [vpn-ipv4] [ipv6] [vpn-ipv6] [mcast-ipv4] [l2-vpn]  
[mvpn-ipv4] [mdt-safi] [ms-pw] [flow-ipv4] [route-target]  
[mcast-vpn-ipv4] [mvpn-ipv6] [flow-ipv6] [evpn] [mcast-ipv6].  
BGP/Group/Neighbor level.

# Routing protocols - Policy Framework: prefix-list

- Prefix-list can contain multiple prefixes
- Address families IPv4, IPv6 or both in 1 prefix-list
- Options for subnet matching
  - Exact
  - Longer
  - Through
  - Range

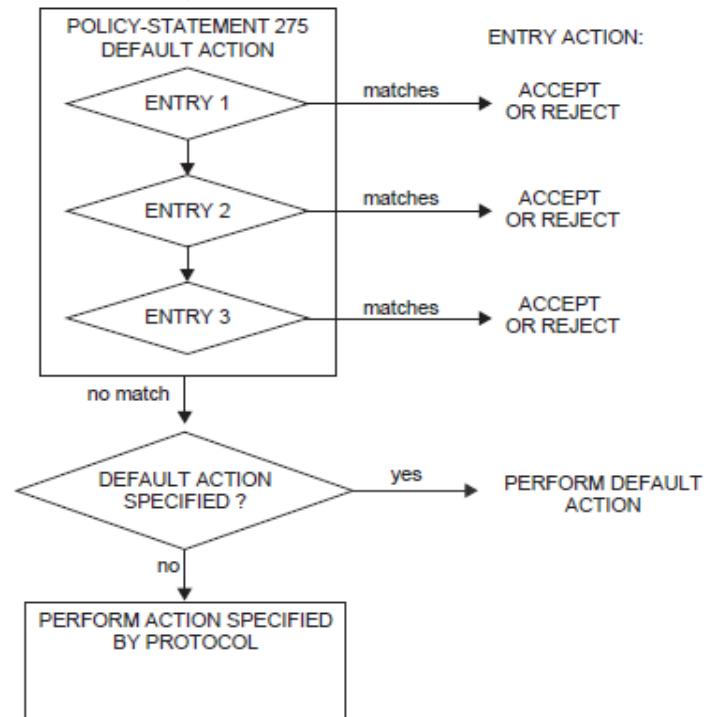
```
configure {
    policy-options prefix-list pref-list-one prefix

    [ip-prefix] (<ipv4-address-and-prefix> | <ipv6-address-and-prefix>)
    <ipv4-address-and-prefix> - <d.d.d.d>(/<d>| <d.d.d.d>)
    <ipv6-address-and-prefix> - (<x:x:x:x:x:x:x:x>|<x:x:x:x:x:x:d.d.d.d>) (/<d>|
    (<x:x:x:x:x:x:x>|<x:x:x:x:x:x:d.d.d.d>))

    IP prefix associated with prefix length
```

# Routing protocols - Policy Framework: policy-statement

```
configure {
    policy-statement "pol-275" {
        entry 1 {
            from {
                community {
                    name "comm-four"
                }
            }
            action {
                action-type accept
                community {
                    add ["comm-five"]
                }
            }
        }
        entry 2 {
            from {
                community {
                    name "comm-seven"
                }
            }
            action {
                action-type reject
            }
        }
        entry 3 {
            from {
                community {
                    name "comm-eight"
                }
            }
            action {
                action-type accept
                as-path {
                    replace "as-one"
                }
            }
        }
        default-action {
            action-type accept
            community {
                add ["comm-six"]
            }
        }
    }
}
```



NOKIA

# Routing protocols - Policy Framework: apply policies

- BGP import and export policies can be applied at BGP, group and neighbor level
- For BGP up to 15 chained policies can be configured (for use with "action next-policy" nesting)

```
configure {
    router "Base" {
        bgp {
            import {
                policy ["pol-276"]
            }
            export {
                policy ["pol-277"]
            }
            group "internal" {
                import {
                    policy ["pol-277" "pol-278"]
                }
                export {
                    policy ["pol-279"]
                }
            }
        neighbor "192.0.2.14" {
            group "internal"
            import {
                policy ["pol-276"]
            }
            export {
                policy ["pol-277"]
            }
        }
    }
}
```

policies can be configured

```
configure {
    router "Base" {
        ospf 0 {
            export-policy ["pol-276"]
        }
    }
}
```

```
configure {
    router "Base" {
        ldp {
            import-policy ["pol-276"]
            export-policy ["pol-277"]
        }
    }
}
```

# EVPN Reduces Operational Complexity

## Provisioning in SROS: as simple as it gets

MON 25 FEB 2019 20:14:00 UTC

[ex:configure service vpls "BD1"]

sap lag-1:16

bgp-evpn evi 16

bgp-evpn mpls auto-bind-tunnel resolution any

bgp-evpn mpls admin-state enable

admin-state enable

[ex:configure service vpls "BD1"]

A:admin@PE-1# info

admin-state enable

service-id 1

customer "1"

bgp-evpn {

    evi 1

    mpls 1 {

        admin-state enable

        auto-bind-tunnel {

            resolution any

        }

    }

    sap lag-1:16 {

    }

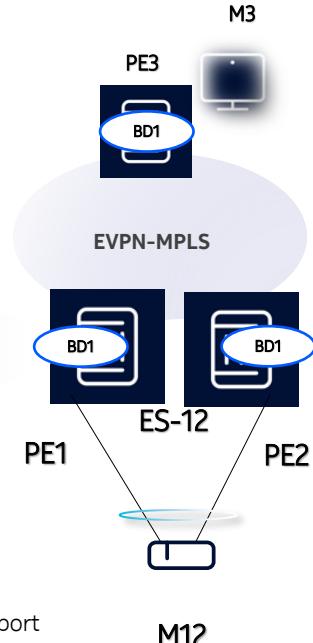
**EVI value <1-65535> provides:  
Auto-derived BGP RD and RT  
Value used for service-carving**

```
A:admin@PE-1# show service id 1 evpn-mpls
=====
BGP EVPN-MPLS Dest
=====
TEP Address      Egr Label    Num. MACs   Mcast      Last Change
Transport:Tnl
-----
192.0.2.3        524251      1           Yes        05/27/2019 08:55:39
          ldp:65538
<snip>
```

```
A:admin@PE-1# show service id 1 fdb detail
=====
Forwarding Database, Service 1
=====
Servid      MAC            Source-Identifier      Type      Last Change
Transport:Tnl-Id
-----
351        00:00:00:00:00:03 eMpls:                Evpn      05/27/19 08:55:39
          192.0.2.3:524252
<snip>
```

**The sap is automatically added  
to the ESI12 based on the port  
on which it is defined**

**SROS/Standard Terminology mapping:**  
SAP or Spoke-SDP == Attachment Circuit  
VPLS == Broadcast Domain (BD)



1

2

3

**Auto-bind-tunnel allows the user to pick up a transport tunnel type:  
ldp, rsvp, bgp, sr-isis, sr-ospf, sr-te, sr-policy, udp, etc.  
'Any' means the best available tunnel will be selected**

# References

The following resources were used to create this lab:

- [AMS-IX Route Servers](#)
- [Implementation of RPKI and IRR filtering on the AMS-IX platform](#)
- [NL-IX Config Guide Peering](#)
- [Nokia SR OS 24.10.1 BGP Guide](#)
- [RFC 7948 & RFC 7948](#)
- [Getting started with BIRD by Kintone](#)
- [HowTo BIRD by dn42](#)
- [IXP Lab by NSRC](#)
- ARouteServer: [repo](#), [docs](#), [tutorial](#)
- [bgpq4](#)
- [RPKI docs by NLNetLabs](#)

# Let's explore the lab together

NOKIA

# Copyright and confidentiality

The contents of this document are proprietary and confidential property of Nokia. This document is provided subject to confidentiality obligations of the applicable agreement(s).

This document is intended for use by Nokia's customers and collaborators only for the purpose for which this document is submitted by Nokia. No part of this document may be reproduced or made available to the public or to any third party in any form or means without the prior written permission of Nokia. This document is to be used by properly trained professional personnel. Any use of the contents in this document is limited strictly to the use(s) specifically created in the applicable agreement(s) under which the document is submitted. The user of this document may voluntarily provide suggestions, comments or other feedback to Nokia in respect of the contents of this document ("Feedback").

Such Feedback may be used in Nokia products and related specifications or other documentation. Accordingly, if the user of this document gives Nokia Feedback on the contents of this document, Nokia may freely use, disclose, reproduce, license, distribute and otherwise commercialize the feedback in any Nokia product, technology, service, specification or other documentation.

Nokia operates a policy of ongoing development. Nokia reserves the right to make changes and improvements to any of the products and/or services described in this document or withdraw this document at any time without prior notice.

The contents of this document are provided "as is". Except as required by applicable law, no warranties of any kind, either express or implied, including, but not limited to, the implied warranties of merchantability and fitness for a particular

purpose, are made in relation to the accuracy, reliability or contents of this document. NOKIA SHALL NOT BE RESPONSIBLE IN ANY EVENT FOR ERRORS IN THIS DOCUMENT or for any loss of data or income or any special, incidental, consequential, indirect or direct damages howsoever caused, that might arise from the use of this document or any contents of this document.

This document and the product(s) it describes are protected by copyright according to the applicable laws.

Nokia is a registered trademark of Nokia Corporation. Other product and company names mentioned herein may be trademarks or trade names of their respective owners.